

Computer Systems 1  
Lecture 12

# Procedures and the Call Stack

Dr. John T. O'Donnell  
School of Computing Science  
University of Glasgow

Copyright ©2019 John T. O'Donnell

# Topics

- 1 Procedures
- 2 Call and return
- 3 Parameter passage
- 4 Procedure calls another procedure
- 5 Saving state
- 6 Stack

# Procedures: reusable code

- Often there is a sequence of instructions that comes up again and again
  - ▶ For example: `sqrt` (square root)
  - ▶ It takes a lot of instructions to calculate a square root
  - ▶ An application program may need a square root in many different places
- We don't want to keep repeating the code
  - ▶ It's tedious
  - ▶ It wastes space (all those instructions require memory!)
- The aim: **write it once** and **reuse the same instructions many times**

# Procedure

- Write the code **one time** — the block of code is called a procedure (or subroutine, function)
- Put the instructions off by themselves somewhere, not in the main flow of instructions
- Give the block of code a label (e.g. work) that describes what it does
- Every time you need to perform this computation, **call it**: go to work
- When it finishes, **the procedure needs to return**: go back to the instruction **after** the one that jumped to it

# Call and return

- One idea is just to use jump instructions for both call and return
- But that isn't actually sufficient — let's look in more detail at what happens

# Returning to the instruction after the call

- Suppose a procedure named **dowork** is used in several places
- Each call jumps to the same place (the address of the first instruction of the procedure)
- But the calls **come from different places**
- Therefore the procedure must finish by **returning to different places**

## Calling and returning

Here is a main program that calls a procedure “dowork” several times. (It takes the value in R1 and doubles it, and the main program would use the result but we ignore that here.)

```

load      R1,x[R0]      ; R1 = x
“goto”    dowork[R0]    ; call the procedure “dowork”
load      R1,y[R0]      ; R1 = y
“goto”    dowork[R0]    ; call dowork
sub       R5,R6,R7      ; R5 = R6-R7

```

```

dowork    add          R1,R1,R1      ; R1 = R1+R1
          “return”      ; return — go back to the caller

```

The return must go to different points in the program—how can it do that?

## Calling and returning

Here is a main program that calls a procedure “dowork” several times. (It takes the value in R1 and doubles it, and the main program would use the result but we ignore that here.)

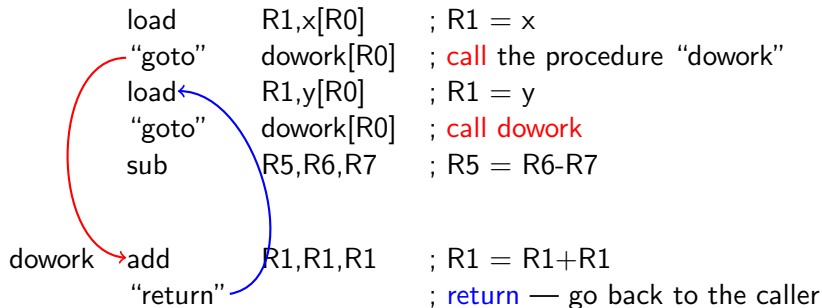
	load	R1,x[R0]	; R1 = x
	“goto”	dowork[R0]	; <b>call</b> the procedure “dowork”
	load	R1,y[R0]	; R1 = y
	“goto”	dowork[R0]	; <b>call dowork</b>
	sub	R5,R6,R7	; R5 = R6-R7
dowork	→ add	R1,R1,R1	; R1 = R1+R1
	“return”		; <b>return</b> — go back to the caller

The return must go to different points in the program—how can it do that?



## Calling and returning

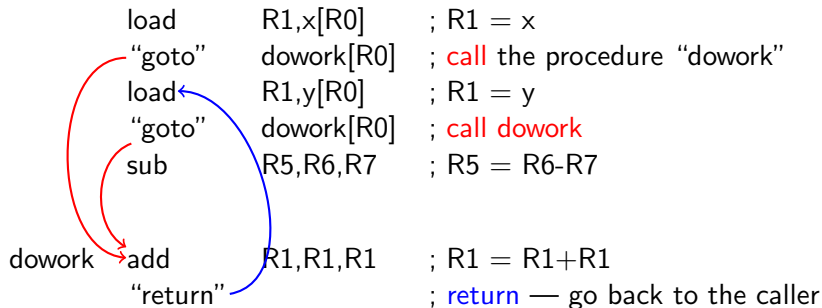
Here is a main program that calls a procedure “dowork” several times. (It takes the value in R1 and doubles it, and the main program would use the result but we ignore that here.)



The return must go to different points in the program—how can it do that?

## Calling and returning

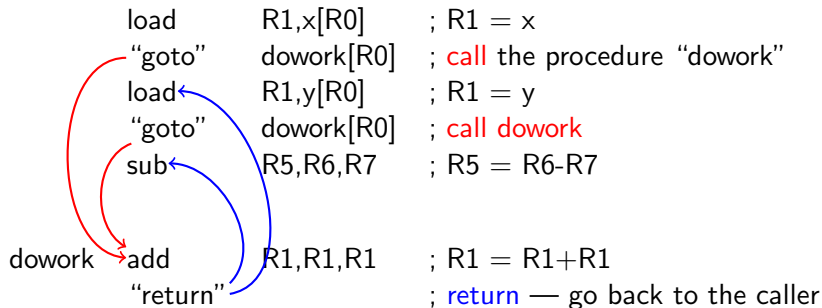
Here is a main program that calls a procedure “dowork” several times. (It takes the value in R1 and doubles it, and the main program would use the result but we ignore that here.)



The return must go to different points in the program—how can it do that?

## Calling and returning

Here is a main program that calls a procedure “dowork” several times. (It takes the value in R1 and doubles it, and the main program would use the result but we ignore that here.)



The return must go to different points in the program—how can it do that?

# The jump-and-link instruction: jal

- When the main program calls the subroutine, it needs to **remember where the call came from**
- This is the purpose of the **jal** instruction — jump and link
- **jal R5,dowork[R0]**
  - ▶ A **pointer to the next instruction after the jal** — the return address — is loaded into the destination register (e.g. R5)
  - ▶ Then the machine jumps to the effective address

# Jumping

All jump instructions (jump, jal, jumplt, etc.) refer to **effective addresses**

- jump **loop**[R0]  
goto loop
- jump **0**[R14]  
goto instruction whose address is in R14
- jump **const**[R2]  
goto instruction whose address is  $\text{const} + R2$

# Implementing call and return

- To call a procedure dowork: `jal R13,dowork[R0]`
  - ▶ The address of the instruction *after* the jal is placed in R13
  - ▶ The program jumps to the effective address, and the procedure starts executing
- To return when the procedure has finished: `jump 0[R13]`
  - ▶ The effective address is  $0 +$  the address of the instruction after the jal
  - ▶ The program jumps there and the main program resumes

# Calling with jal and returning with jump

```

load    R1,x[R0]           ; R1 = x
jal     R13,dowork[R0]     ; call dowork
load    R1,y[R0]           ; R1 = y
jal     R13,dowork[R0]     ; call dowork
sub     R5,R6,R7           ; R5 = R6-R7

```


```

dowork  add    R1,R1,R1     ; R1 = R1+R1
        jump   0[R13]      ; return

```

- **call** — jal puts a pointer to the next instruction into R13
- **return** — follow the pointer in R13

# Calling with jal and returning with jump



```

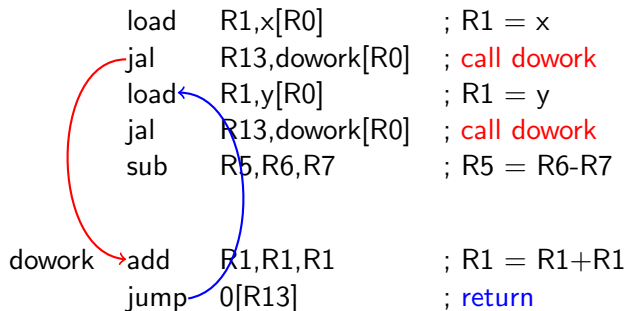
    load    R1,x[R0]           ; R1 = x
    jal     R13,dowork[R0]     ; call dowork
    load    R1,y[R0]           ; R1 = y
    jal     R13,dowork[R0]     ; call dowork
    sub     R5,R6,R7           ; R5 = R6-R7

dowork: add   R1,R1,R1           ; R1 = R1+R1
        jump 0[R13]             ; return
  
```

- **call** — jal puts a pointer to the next instruction into R13
- **return** — follow the pointer in R13

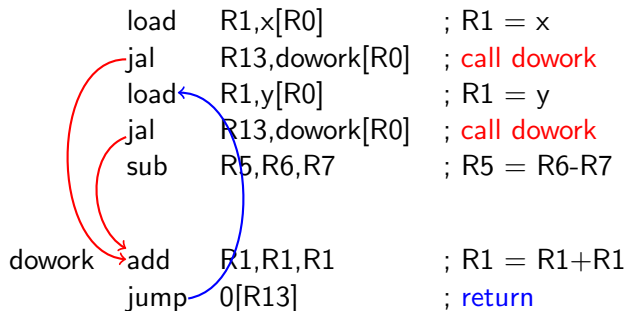


# Calling with jal and returning with jump



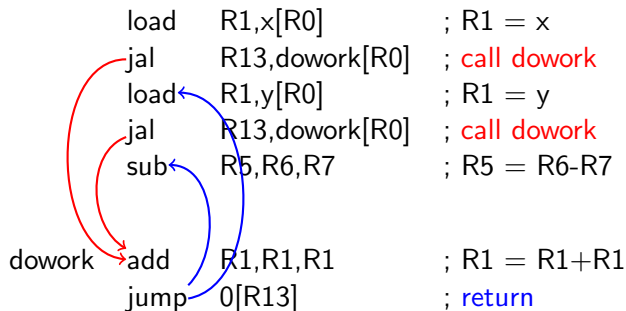
- **call** — jal puts a pointer to the next instruction into R13
- **return** — follow the pointer in R13

# Calling with jal and returning with jump



- **call** — jal puts a pointer to the next instruction into R13
- **return** — follow the pointer in R13

# Calling with jal and returning with jump



- **call** — jal puts a pointer to the next instruction into R13
- **return** — follow the pointer in R13

# Parameter passage

- There are several different conventions for passing argument to the function, and passing the result back
- What is important is that the caller and the procedure agree on how information is passed between them
- If there is a small number of arguments, the caller may put them in registers before calling the procedure
- If there are many arguments, the caller builds an array or vector (sequence of adjacent memory locations), puts the arguments into the vector, and passes the address of the vector in a register (typically R1)
- A simple convention: **the argument and result are passed in R1**

# Functions

- A **function** is a procedure that
  - ▶ Receives a parameter (a word of data) from the caller
  - ▶ Calculates a result
  - ▶ Passes the result back to the caller when it returns
- A **pure function** is a function that doesn't do anything else — it doesn't change any global variables, or do any input/output

## Example: Passing argument and result in R1

; Main program

```

load  R1,x[R0]          ; arg = x
jal   R13,work[R0]      ; result = work (x)
...
load  R1,y[R0]          ; arg = y
jal   R13,work[R0]      ; result = work (y)
...
```

; Function work (x) = 1 + 7\*x

```

work  lea    R2,7[R0]    ; R7 = 2
      lea    R3,1[R0]    ; R3 = 1
      mul    R1,R1,R2     ; result = arg * 7
      add    R1,R3,R1     ; result = 1 + 7*arg
      jump   0[R13]      ; return
```

# What if a procedure calls another procedure?

- The simplest kind of procedure
  - ▶ Call it with `jal R13,procname[R0]`
  - ▶ It returns by executing `jump 0[R13]`

# Limitations of basic call

- If the procedure modifies any registers, it may destroy data belonging to the caller
- If the procedure calls another procedure, it can't use R13 again. Each procedure would need a dedicated register for its return address, limiting the program to a small number of procedures
- The basic call mechanism doesn't allow a procedure to call itself (this is called **recursion**)



# R13 overwritten: proc1 returns to the wrong place!

```

load    R1,x[R0]          ; R1 := x
jal     R13,proc1[R0]      ; call proc1
load    R1,y[R0]          ; R1 := y

```

R13

```

proc1   store    R1,result[R0] ; result := 2 × (2x)2
proc1   lea      R2,1[R0]      ; R2 := 1
proc1   add      R1,R1,R1      ; R1 := R1+1
proc1   jal      R13,proc2[R0] ; R1 := square (R1)
proc1   add      R1,R1,R1      ; R1 := R1+1
proc1   jump     0[R13]        ; return

```

```

square  mul      R1,R1,R1      ; R1 = R1*R1
square  jump     0[R13]        ; return

```

# R13 overwritten: proc1 returns to the wrong place!

```

load    R1,x[R0]          ; R1 := x
jal     R13,proc1[R0]      ; call proc1
load    R1,y[R0]          ; R1 := y

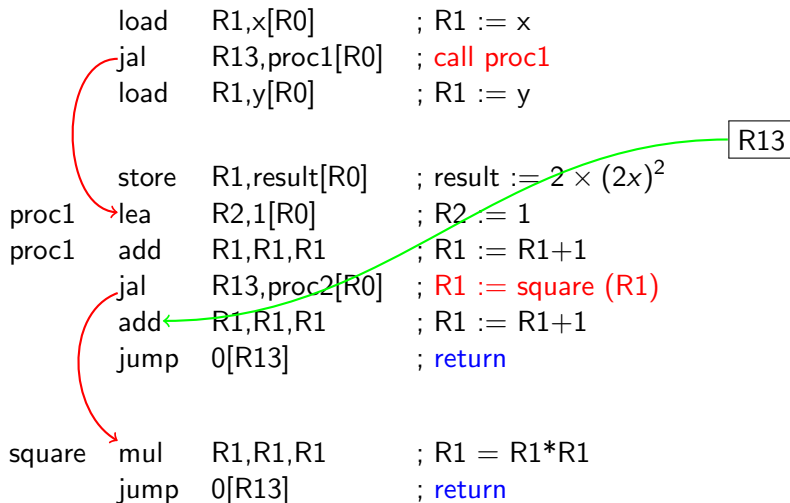
proc1   store R1,result[R0] ; result := 2 × (2x)2
proc1   lea   R2,1[R0]      ; R2 := 1
proc1   add   R1,R1,R1      ; R1 := R1+1
proc1   jal   R13,proc2[R0] ; R1 := square (R1)
proc1   add   R1,R1,R1      ; R1 := R1+1
proc1   jump  0[R13]        ; return

square  mul   R1,R1,R1      ; R1 = R1*R1
square  jump  0[R13]        ; return

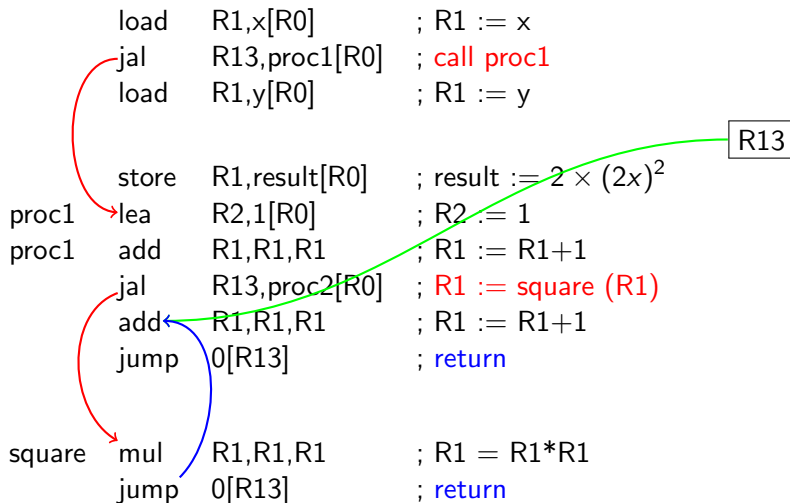
```

The diagram illustrates a bug in the assembly code. A red arrow points from the `jal R13,proc1[R0]` instruction to the `proc1` label, indicating a jump to the start of the `proc1` procedure. A green arrow points from the `load R1,y[R0]` instruction to the `R13` register box, indicating that the value of `R13` is being overwritten. The `proc1` procedure then jumps to `0[R13]`, which is the address stored in `R13` (the address of the `load R1,y[R0]` instruction), causing it to return to the wrong place.

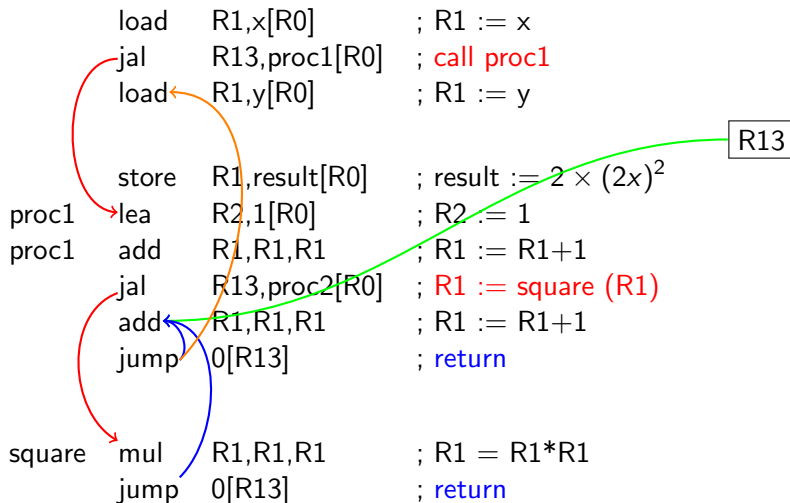
# R13 overwritten: proc1 returns to the wrong place!



## R13 overwritten: proc1 returns to the wrong place!



# R13 overwritten: proc1 returns to the wrong place!



# Saving state

- Calling a procedure creates new information
  - ▶ The return address
  - ▶ Whatever values the procedure loads into the registers
- But this new information could overwrite essential information belonging to the caller
- We need to **save the caller's state** so the procedure won't destroy it

# The wrong way to save state

- Suppose we just have a variable saveRetAdr
- Store R13 into it in the procedure, load that when we return
- Now it's ok for proc1 to call proc2
- But if proc2 calls proc3 we are back to the same problem: it doesn't work!
- The solution: a **stack**

# Saving registers

- Most procedures need to use several registers
- It's nearly impossible to do *anything* without using some registers!
- The first thing a procedure should do is to *save the registers* it will use by copying them into memory (with store instructions).
- The last thing it should do before returning is to *restore the registers* by copying their values back from memory (with load instructions).



# Where can the registers be saved?

- It won't work to copy data from some of the registers to other registers!
- It's essential to save the data into memory
- Two approaches
  - ▶ Allocate fixed variables in memory to save the registers into — simple but doesn't allow recursion
  - ▶ Maintain a **stack** in memory, and **push** the data onto the stack — this is the best approach and is used by most programming languages

# Who saves the state: the caller or the procedure?

- Two approaches:
- Caller saves (used occasionally)
  - ▶ Before calling a procedure, the caller saves the registers, so all its essential data is in memory
  - ▶ After the procedure returns, the caller does whatever loads are needed
- Callee saves (usually the preferred solution)
  - ▶ The caller keeps data in registers, and assumes that the procedure won't disturb it
  - ▶ The first thing the procedure does is to save the registers it needs to use into memory
  - ▶ Just before returning, the procedure restores the registers by loading the data from memory

# Stack of return addresses

- To allow a large number of procedures, we can't dedicate a specific register to each one for its return address
- Therefore we
  - ▶ Always use the same register for the return address in a jal instruction (we will use R13)
  - ▶ The first thing a procedure does is to store its return address into memory
  - ▶ The last thing the procedure does is to load its return address and jump to it
  - ▶ The return addresses are pushed onto a *stack*, rather than being stored at a fixed address


# Stacks

- A **stack** is a container
- Initially it is empty
- You can **push** a value onto the stack; this is now sitting on the top of the stack
- You can **pop** the stack; this removes the **most recently pushed** value and returns it
- A stack allows access only to the top value; you cannot access anything below the top
- We can save procedure return addresses on a stack because return always needs the most recently saved return address


Initially the stack is empty




# Call procedure, push return address $a$



Call another procedure, push return address  $b$




Return: pop produces return address  $b$





Call some procedure, push return address *c*



Call a procedure, push return address  $d$

$d$
$c$
$a$

# The call stack

- Central technique for
  - ▶ Preserving data during a procedure call
  - ▶ Holding most of your variables
- It goes by several names; these are all the same thing
  - ▶ call stack
  - ▶ execution stack
  - ▶ “The stack”
- It's important!
  - ▶ Most programming languages use it
  - ▶ Computers are designed to support it
  - ▶ Often referred to (Stack Overflow web site, etc.)

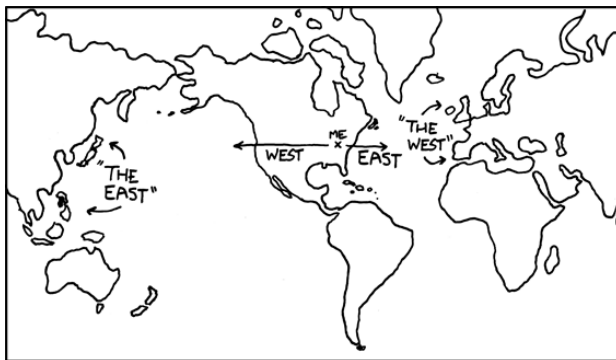
# Stack frames

- There is a **call stack** or **execution stack** that maintains complete information about all procedure calls and returns
- Every “activation” of a procedure pushes a **stack frame**
- When the procedure returns, its stack frame is popped (removed) from the stack
- R14 contains the address of the current (top) stack frame
- The stack frame contains:
  - ▶ A pointer to the previous stack frame (this is required to make the pop work)
  - ▶ The return address (saved value of R13)
  - ▶ The saved registers (so the procedure can use the registers without destroying information)
  - ▶ Local variables (so the procedure can have some memory of its own to use)

# Implementing the call stack

- Dedicate R14 to the **stack pointer**
- This is a programming convention, not a hardware feature
- When the program is started, R14 will be set to point to an empty stack
- When a procedure is called, the saved state will be pushed onto the stack: store a word at  $0[R14]$  and add 1
- When a procedure returns, it pops the stack and restores the state: subtract 1, load from  $0[R14]$
- The program should never modify R14 apart from the push and pop

## terminology



THIS ALWAYS BUGGED ME.

<https://xkcd.com/503/>