

Computer Systems 1  
Lecture 16

# Nested Conditionals

Dr. John T. O'Donnell  
School of Computing Science  
University of Glasgow

Copyright ©2019 John T. O'Donnell

# Topics

- 1 Blocks
  - Syntax for blocks
  - Enter at beginning, exit at end
- 2 Nested conditionals
- 3 Case and jump tables
- 4 Real world incident: the SSL Bug

# Blocks

- You'll need to learn many programming languages, eventually
- There are concepts that appear in most languages
- It's useful to focus on the general concept
- One important concept is a block
- A **block** or **compound statement** is a single statement that contains several statements
- The purpose is to let you have a group of statements in a loop, or controlled by a conditional
- The **syntax** is the detailed punctuation used to indicate a block, and this varies in different languages
  - ▶ There are lots of small syntax differences between languages
  - ▶ Some languages use `:=` to mean assign, `=` to mean equals
  - ▶ Some languages use `=` to mean assign, `==` to mean equals

# Python syntax style for blocks: layout

The layout (the indentation) determines what is inside the if statement

```
a = 1
if x < y :
    b = 2
    c = 3
d = 4
```

If you change the indentation, you change the meaning of the program

In Python, you write `:` to mean *then*, and `else` to mean *else*

## Algol/Pascal style for blocks: begin—end

```
a := 1;
if x < y
  then begin b := 2;
           c := 3
        end
  else begin d := 4;
           e := 5
        end;
f := 6;
```

You write `then` to mean *then*, and `else` to mean *else*

Statements must be separated by semicolon ;

## C style for blocks: braces

```
a = 1;  
if (x < y) b = 2;  
else  
    { d = 4;  
      e = 5; }  
f = 6;
```

In C you don't write `then` at all, but this means the condition `x < y` must be enclosed in parentheses so the compiler can tell where the condition ends and the then-statement begins

## Block structured style for blocks: matching keywords

```
a := 1;
if x < y
  then b := 2;
       while i < n do
           sum := sum + x[i];
           i := i + 1
       endwhile
  else d := 4
       e := 5
endif;
f := 6;
```

This style introduces a lot of keywords (endif, endwhile, endfor, endrepeat) but you don't need braces around a block. It makes code more readable and enables the compiler to produce better error messages

# Enter at beginning, exit at end

- A common programming style is to require
  - ▶ Each block enters only at the beginning of the block
  - ▶ Each block exits only at the end of the block
- This style is helpful in some programming languages, but in some languages it sometimes makes code less readable
- In assembly language, and for compilation patterns, it is necessary to follow this style
- In high level languages this style is sometimes helpful, but not always



# Single entrance/exit for compilation patterns

- It is straightforward to translate high level control constructs using the compilation patterns
- These *require* that the blocks of code always start at the beginning and finish at the end
- It's bad to jump into the middle of a block, or exit out of the middle because
  - ▶ The compilation patterns won't work correctly
  - ▶ You'll have to duplicate a lot of code
  - ▶ Example: returning from a procedure requires restoring the registers, resetting the stack pointer, loading the return address
  - ▶ That code should not be duplicated in several places in a procedure

# Systematic approach to programming?

- Start by understanding what your program should do
- Express the algorithm using high level language notation (and it's ok to mix in some English too)
- Translate the high level to the low level
  - ▶ Assignment statements:  $x := \text{expression}$
  - ▶ I/O statements: Write string
  - ▶ goto L
  - ▶ if boolean then goto L
- Translate the low level to assembly language
- Retain the high and low level code as comments
- Do hand execution at every level

# Why use this systematic approach to programming?

- It enables you to write correct code at the outset, and minimise debugging
- If there is a bug, it helps you to catch it early (e.g. in translation to goto form)
- If there's a bug in an instruction, the comments enable you to find it quickly
  - ▶ A common error is to use a wrong register number: `add R9,R3,R4`
  - ▶ Poor comments don't help: `add R9,R3,R4 ; R9 := R3+R4`
  - ▶ Good comments help a lot: `add R9,R3,R4 ; x := alpha + (a[i]*b[i])`
  - ▶ Look at the register usage comments (oops, x is in R8, not R9, now I know how to fix it and I don't have to read the entire program)
- Professional software needs to be maintained; the comments make the software easier to read and more valuable

# Nested if-then-else

Conditional statements can be nested deeply

```
if b1
  then S1
      if b2
          then S3
          else S4
      S5
  else S6
      if b3
          then S7
      S8
```

## Special case of nested if-then-else

Often the nesting isn't random, but has this specific structure:

```
if b1
  then S1
  else if b2
    then S2
    else if b3
      then S3
      else if b4
        then S4
        else if b5
          then S5
          else S6
```

(Some languages require punctuation to avoid ambiguity)

## Another way to write it

To avoid running off the right side of the window, it's usually indented like this:

```
if b1
  then S1
else if b2
  then S2
else if b3
  then S3
else if b4
  then S4
else if b5
  then S5
else S6
```

## Some programming languages have elsif or elif

```
if b1 then S1
  elif b2 then S2
  elif b3 then S3
  elif b4 then S4
  elseif b5 then S5
  else S6
```

- It avoids ambiguity
- It signals to the compiler and to a human programmer that this specific construct is being used
- It allows good indentation layout without violating the basic principle of indentation
- Some languages have this, some don't

## A common application: numeric code

Nested if-then-else but the boolean conditions are not arbitrary: they are checking the value of a code:

```
if code = 0
  then S1
else if code = 1
  then S2
else if code = 2
  then S3
else if code = 3
  then S4
else if code = 4
  then S5
```



# The case statement

```
case n of
  0 -> Stmt
  1 -> Stmt
  2 -> Stmt
  3 -> Stmt
  4 -> Stmt
  5 -> Stmt
  else -> Stmt // handle error
```

This means: execute the statement corresponding to the value of  $n$

Many programming languages have this; the syntax varies a lot but that isn't what's important

## Example: numeric code specifies a command

```
; The input data is an array of records, each specifying an operation
;   Command : record
;       code : Int      ; specify which operation to perform
;       i    : Int      ; index into array of lists
;       x    : Int      ; value of list element

; The meaning of a command depends on the code:
;   0  terminate the program
;   1  insert x into set[i]
;   2  delete x from set[i]
;   3  return 1 if set[i] contains x, otherwise 0
;   4  print the elements of set[i]
```

# Selecting the command with a case statement

```

; Initialize
;   BuildHeap ()

; Execute the commands in the input data
;   finished := 0
;   while InputPtr <= InputEnd && not finished
;       CurrentCode := (*InputPtr).code
;       p := set[*InputPtr]      ; linked list
;       x := (*InputPtr).x      ; value to insert, delete, search

;       case CurrentCode of
;           0 : <CmdTerminate>
;           1 : <CmdInsert>
;           2 : <CmdDelete>
;           3 : <CmdSearch>
;           4 : <CmdPrint>
;           else : <>
;       InputPtr := InputPtr + sizeof(Command)
; Terminate the program
;   halt

```

# Finding a numeric code

- It's tedious and inefficient to go through the possible values of a numeric code in sequence
  - ▶ If you're looking up Dr Zhivago in the phone book, do you look at Arnold Aardvark, and Anne Anderson, on on and on?
  - ▶ You go straight to the end of the book
- We want to find the statement corresponding to a numeric code directly, without checking all the other values

# A problem with efficiency

- The problem
  - ▶ There are many applications of case statements
  - ▶ They are executed frequently
  - ▶ The number of cases can be large (not just 4 or 5; can be hundreds)
  - ▶ The implementation of the if-then-else requires a separate compare and jump for each condition
- The solution
  - ▶ A technique called **jump tables**

# Jump tables: the basic idea

- For each target statement ( $S_1$ ,  $S_2$ ,  $S_3$ , etc) in the conditional, introduce a jump to it: `jump S1[R0]`, `jump S2[R0]`, etc
- Make an **array** “jt” of these jump instructions

`jt[0] = jump S0[R0]`

`jt[1] = jump S1[R0]`

`jt[2] = jump S2[R0]`

`jt[3] = jump S3[R0]`

`jt[4] = jump S4[R0]`

- Given the *code*, Jump to `jt[code]`
- Each element of the array is an instruction that requires two words
- So jump to  $\&\text{jt} + 2 \times \text{code}$

# Jump table

```
; Jump to operation specified by code
    add    R4,R4,R4           ; code := 2*code
    lea    R5,CmdJumpTable[R0] ; R5 := pointer to jump table
    add    R4,R5,R4           ; address of jump to operation
    jump   0[R4]              ; jump to jump to operation

CmdJumpTable
    jump   CmdTerminate[R0]   ; code 0: terminate the program
    jump   CmdInsert[R0]      ; code 1: insert
    jump   CmdDelete[R0]      ; code 2: delete
    jump   CmdSearch[R0]      ; code 3: search
    jump   CmdPrint[R0]       ; code 4: print

CmdDone
    load   R5,InputPtr[R0]
    lea    R6,3[R0]
    add    R5,R5,R6
    store  R5,InputPtr[R0]
    jump   CommandLoop[R0]
```

# We have to be careful!

- What if code is negative, or larger than the number of cases?
- The jump to the jump table could go anywhere!
- It might not even go to an instruction
- But whatever is in memory at the place it goes to, will be interpreted as an instruction that will be executed
- The program will go haywire
- Debugging it will be hard: the only way to catch the error is to read the code and/or to single step
- Solution: before jumping into the jump table, check to see if code is invalid (too big or too small)



# Checking whether the code is invalid

## CommandLoop

```

    load   R5,InputPtr[R0]      ; R1 := InputPtr
    load   R4,0[R5]             ; R4 := *InputPtr.code
; Check for invalid code
    cmp    R4,R0                ; compare (*InputPtr).code, 0
    jumplt CmdDone[R0]          ; skip invalid code (negative)
    lea    R6,4[R0]             ; maximum valid code
    cmp    R4,R6                ; compare code with max valid code
    jumpgt CmdDone[R0]          ; skip invalid code (too large)

```

...

## CmdDone

```

    load   R5,InputPtr[R0]
    lea    R6,3[R0]
    add    R5,R5,R6
    store  R5,InputPtr[R0]
    jump   CommandLoop[R0]

```

# A real world incident!

The moral of the following true story: be meticulous!

- On 21 February 2014, Apple announced a bug in core security software
- It affected iPhone, iPad, Macintosh, Safari
- The bug is serious: it skips a crucial security check in SSL
- A security update was made available quickly
  - ▶ That's why you should always apply your security updates!
- It was fixed quickly but we can learn some important lessons from it

# What is SSL?

- Secure Sockets Layer
- It arranges that communication between your device and a server is **encrypted**
- When you see a “padlock icon” on a browser, it means the SSL protocol is being used

## Source file: sslKeyExchange.c

- 1,969 lines of code in C
- You can see the code—click on this link, or type it in:
- `http://opensource.apple.com/source/Security/Security-55471/libsecurity\_ssl/lib/sslKeyExchange.c`
- The bug made it possible to break open “secure” Internet traffic
- It appears somewhere on the next three slides: can you spot it?

# The function SSLVerifySignedServerKeyExchange (1)

```
static OSStatus
SSLVerifySignedServerKeyExchange(SSLContext *ctx, bool isRsa, SSLBuffer signedParams,
                                uint8_t *signature, UInt16 signatureLen)
{
    OSStatus      err;
    SSLBuffer      hashOut, hashCtx, clientRandom, serverRandom;
    uint8_t        hashes[SSL_SHA1_DIGEST_LEN + SSL_MD5_DIGEST_LEN];
    SSLBuffer      signedHashes;
    uint8_t        *dataToSign;
    size_t         dataToSignLen;

    signedHashes.data = 0;
    hashCtx.data = 0;

    clientRandom.data = ctx->clientRandom;
    clientRandom.length = SSL_CLIENT_SRVR_RAND_SIZE;
    serverRandom.data = ctx->serverRandom;
    serverRandom.length = SSL_CLIENT_SRVR_RAND_SIZE;
```

# The function continued (2)

```

if(isRsa) {
    /* skip this if signing with DSA */
    dataToSign = hashes;
    dataToSignLen = SSL_SHA1_DIGEST_LEN + SSL_MD5_DIGEST_LEN;
    hashOut.data = hashes;
    hashOut.length = SSL_MD5_DIGEST_LEN;

    if ((err = ReadyHash(&SSLHashMD5, &hashCtx)) != 0)
        goto fail;
    if ((err = SSLHashMD5.update(&hashCtx, &clientRandom)) != 0)
        goto fail;
    if ((err = SSLHashMD5.update(&hashCtx, &serverRandom)) != 0)
        goto fail;
    if ((err = SSLHashMD5.update(&hashCtx, &signedParams)) != 0)
        goto fail;
    if ((err = SSLHashMD5.final(&hashCtx, &hashOut)) != 0)
        goto fail;
}
else {
    /* DSA, ECDSA - just use the SHA1 hash */
    dataToSign = &hashes[SSL_MD5_DIGEST_LEN];
    dataToSignLen = SSL_SHA1_DIGEST_LEN;
}
hashOut.data = hashes + SSL_MD5_DIGEST_LEN;
hashOut.length = SSL_SHA1_DIGEST_LEN;

```

# The function continued (3)

```

if ((err = SSLFreeBuffer(&hashCtx)) != 0)
    goto fail;

if ((err = ReadyHash(&SSLHashSHA1, &hashCtx)) != 0)
    goto fail;
if ((err = SSLHashSHA1.update(&hashCtx, &clientRandom)) != 0)
    goto fail;
if ((err = SSLHashSHA1.update(&hashCtx, &serverRandom)) != 0)
    goto fail;
if ((err = SSLHashSHA1.update(&hashCtx, &signedParams)) != 0)
    goto fail;
if ((err = SSLHashSHA1.final(&hashCtx, &hashOut)) != 0)
    goto fail;

err = sslRawVerify(ctx,
                   ctx->peerPubKey,
                   dataToSign,
                   dataToSignLen,
                   signature,
                   signatureLen);
/* plaintext */
/* plaintext length */

if(err) {
    sslErrorLog("SSLDecodeSignedServerKeyExchange: sslRawVerify "
               "returned %d\n", (int)err);
    goto fail;
}

fail:
    SSLFreeBuffer(&signedHashes);
    SSLFreeBuffer(&hashCtx);
    return err;
}

```

# A bit of C

## Syntax of the if-statement

```
if (condition) statement;
```

Examples:

```
if (x<y) n=n+1;
```

```
if (n==5)
{
    p = 3;
    q = 4;
    printf ("hello p = %3d\n", p);
}
```



## A closer look

```
if ((err = SSLHashSHA1.update(&hashCtx, &serverRandom)) != 0)
    goto fail;
if ((err = SSLHashSHA1.update(&hashCtx, &signedParams)) != 0)
    goto fail;
    goto fail;
if ((err = SSLHashSHA1.final(&hashCtx, &hashOut)) != 0)
    goto fail;
```

The second goto fail is **unconditional** so the check in the last if statement is **never executed**

# What can we learn from this?

- To be a good programmer you need to
  - ▶ Know your programming language
  - ▶ Be meticulous
  - ▶ **Read** your code carefully
  - ▶ Do thorough testing
- It was very helpful that Apple published the source code for this bug

# A mystery: The Attempted Linux Trapdoor Hack of 2003

- A **trapdoor** is a flaw (bug? hack?) in software that allows a hostile user to gain abilities they should not have
  - ▶ It's like giving the key to the safe in a bank to a burglar
- The master copy of Linux was maintained on Bitkeeper (no longer—now git is used)
- A secondary copy was on cvs
- Someone—unknown to this day—made a small edit to an obscure piece of code on the cvs version

## The offending code: can you spot the trapdoor?

- A change was made to a function that a user could call to wait for an event
- System functions typically take options, and they need to check to see if the options are valid
- If not, a return code `retval` is set to a non-zero value

So, here's the code:

```
if ((options == (__WCLONE|__WALL)) && (current->uid = 0))  
    retval = -EINVAL;
```

# What's going on?

- In C, `==` is the comparison operator
- `=` is the assignment operator (bad notation, but we're stuck with it)
- This code looks like it is checking:
  - ▶ If either the `WCLONE` or `WALL` option flag is set and the current user is root, then return an error condition
- What it **actually** says is
  - ▶ Check the irrelevant flags, **set the current user to be root**, get the irrelevant result 0 and do nothing else


```
if ((options == (__WCLONE|__WALL)) && (current->uid = 0))
    retval = -EINVAL;
]
```

# What's the significance of root?


- Many **processes** are running in the computer
- Each has its own set of privileges
- There is one special user “root” which has *all* privileges: **root can do anything at all**
- The effect of this change to the source code (a change of *one character!* is
  - ▶ If a user program calls this obscure function with just the right set of obscure options, **it suddenly gains full control over the machine**
- Don't worry — **this faulty code never made it into the master copy of Linux**

## HIGHWAY ENGINEER PRANKS:

## THE INESCAPABLE CLOVERLEAF:



## THE ZERO-CHOICE INTERCHANGE:



## THE ROTARY SUPERCOLLIDER:



<https://xkcd.com/253/>