

# Computer Systems 1

## Lecture 17

# Trees

Dr. John T. O'Donnell  
School of Computing Science  
University of Glasgow

Copyright ©2019 John T. O'Donnell

# Topics

- 1 Survey about programming in Python
- 2 Assessed exercise: ordered lists program
- 3 Trees

# Survey about programming in Python

- This is completely optional
- There is a research group investigating the process of learning a programming language
- We are running two surveys
  - ① Starting Tuesday 5 March at 1pm, closing Tuesday 12 March at 12:00 noon
  - ② The lecture on Tuesday 12 March will be about programming language semantics: relating Python to compilation patterns and machine language
  - ③ Followup survey starting Tuesday 12 March at 1pm, closing Tuesday 19 March at 12:00 noon
  - ④ The last lecture of the course is Friday 21 March, and will discuss the results
- Participation is **optional** and **anonymous**.
- It will not affect your grade in any way, but we hope you find it interesting and helpful in learning programming languages

# Assessed exercise: ordered lists program

- ① There is one lab exercise that will be assessed: it counts for 10% of your grade in the course
- ② It will be posted tonight on Moodle
- ③ You are given a reasonably long program, which contains a few small missing pieces
- ④ The exercise is to
  - ① Read and understand the program
  - ② Complete the missing pieces

# Concepts used in the program

## ① Array of records

- ① Representing a command as a records
- ② Traversing an array of records
- ③ Case statement and jump table

## ② Linked lists

- ① Traversing a list to print its elements
- ② Insertion in list keeping the elements in ascending order
- ③ Deletion from a list
- ④ Searching a list

# Ordered lists

There is an array of lists, initially empty. There are  $n$  lists of them.

```
list[0] = [ ]  
list[1] = [ ]  
...  
list[nlists-1] = [ ]
```

At all times as the program runs, the lists are ordered: their elements are increasing

```
list[0] = [4, 9, 23, 51 ]  
list[1] = [7, 102, 238 ]  
...  
list[nlists-1] = [2, 87, 89, 93, 103, 195 ]
```

# Commands

The program executes commands:

- Terminate — the program finishes
- Insert into list  $i$  the value  $x$  — modify  $\text{list}[i]$  so it contains  $x$ , while maintaining the ascending order
- Delete from list  $i$  the value  $x$  — modify  $\text{list}[i]$  so  $x$  is removed, but don't do anything if  $x$  isn't in the list
- Search list  $i$  for  $x$  — print Yes if  $x$  is in the list, No otherwise
- Print  $i$  — the numbers in  $\text{list}[i]$  are printed

# Example

- Insert into list[3] the value 23 [23]
- Insert into list[3] the value 6 [6, 23]
- Insert into list[3] the value 67 [6, 23, 67]
- Insert into list[3] the value 19 [6, 19, 23, 67]
- Print list[3] 6 19 23 67



# Why are ordered lists useful?

- This is one way to arrange a database: think of the elements as persons' names, or matriculation numbers
- Sometimes you want to process all the data in a container in a specified order
- If the data is ordered, it's faster to find a particular item (on average you only have to check half of the items)
- An ordered list can be used to represent a set

# Where do the commands come from?

- In a real application, we would read the commands from input
- But in this program, each command is represented as a record
- The entire input is a static array of records defined with data statements
- This is easier because
  - ▶ If you read from an input device, it's necessary to convert the input character string to numbers
  - ▶ In testing a program, it's convenient to have input data that is **fixed and repeatable**
  - ▶ Don't want to have to type in the same input every time you run the program!

# Representing a command

- Each command is a record with three fields
  - ▶ A code indicating which kind of command
  - ▶ A number  $i$  indicating which list we're operating on
  - ▶ A value  $x$  which might be inserted etc
- Each record must have these three fields
- Some commands don't use them all (e.g. Print just needs  $i$ , not  $x$ )
- The main program uses a **case** statement to handle each command, and implements this with a **jump table**

# Reading a program before writing

- You should *read and understand* the program before modifying it
  - ▶ Reading a program is an important skill you will need throughout your career
  - ▶ The program is filled with examples so it is excellent revision material
  - ▶ You need to understand a program before you'll be able to make changes to it
- One of the aims of the exercise is to get experience with reading a longer program—don't skip this!

# Some tips on testing and debugging

- Debugging has two phases:
  - 1 Diagnosis: finding out what went wrong and why
  - 2 Correction: fixing the error
- The most important point: don't just make random changes to the code and hope for the best—instead, find out what the error is and fix it cleanly

# Reading and testing a program

- A good way to understand a section of assembly language instructions is to step through it, one instruction at a time
  - ① Check that the instruction did what you expected it to do
  - ② Check that the instruction is consistent with its comment
  - ③ Try to relate the instruction with the bigger picture: what is it doing in the context of the program?
- Coverage
  - ① You don't need to step through a set of instructions a huge number of times
  - ② If there's a loop, step through two or three iterations
  - ③ If possible, arrange test data so the loop will terminate after just a few iterations
  - ④ But try to step through as much of the program as possible
  - ⑤ This is called **coverage**: try to cover all of the program with your testing

# Breakpoints

- It's a good idea to step through a program one instruction at a time, so you understand clearly what each instruction is doing
- However, in a longer program this isn't always feasible
  - ▶ The OrderedLists program has to build the heap when it starts; this may take several thousand instructions before it even really gets going!
- Solution: **breakpoints**
  - ▶ Find the address of an instruction where you want to start single stepping
  - ▶ Enter this address as a breakpoint
  - ▶ Click Run to execute the program at full speed; when it reaches the breakpoint it will stop
  - ▶ Then you can single step to examine what the instructions are doing

# How to set a breakpoint

- On the Processor pane, click Breakpoint. It will say “Breakpoint is off”
- Enter the breakpoint command and click Set Breakpoint
- BPeq BPpc (BPhex "01a6")
- It will say “Breakpoint is on”. Click Close
- On Processor, click Run. It will stop when the pc register gets the value you specified



# Tree

- A node doesn't have to have two fields named *value* and *next* — it's normal to define a specific node type for an application program.
- Nodes with *value* and *next* can be connected into a **linked list**.
- Nodes can also have with several fields containing data, not just one “value” field.
- And a node can have several pointer fields...
- Common case: a **binary tree** has two pointers in each node, named **left** and **right**.
- Each of these can either contain nil, or point to another node.


# Tree

Node : record

```
value      ; the actual data in the node
left       ; left subtree is a pointer to a Node
right      ; right subtree is a pointer to a Node
```

- Similar to a node for a linked list, but with two pointers
- There can also be several fields for data, not just one “value” field
- And we could have more than just two pointers

# A binary tree



In computer science, for some reason we draw trees upside down

Suppose  $p$  is a pointer to the tree

- $(*p).left$  is the pointer to the left subtree
- $(*p).right$  is the pointer to the right subtree

# Applications of trees

Trees are used everywhere in programming

- To hold structured data
- To make programs faster (*much* faster)


# Holding structured data

- A compiler reads in program text, which is just a character string: a sequence of characters.
- It needs to represent the deep structure underlying that sequence of characters.
- This is done by building a tree (the part of a compiler that takes a character string and produces a tree is called the **parser**).

# Parsing

$x := a + b * c$

assignment



## Another application of jump tables!

- In complicated applications, trees normally have **several different types of node**
- Examples: operations with 1 operand; operations with 2 operands; control constructs with a boolean expression and two statements, etc.
- So there are several different kinds of record
- Each record has a *code* in the first word
- The value of the code determines how many more words there are in the record, and what they mean
- When a program has a pointer to a node, it needs to examine the code and take different actions depending on what the code is
- This is done with a jump table

# Searching


- Suppose we have a large number of records (e.g. a database)
- We want to search the database for an entry where a field has a certain value (e.g. search for a record where the MatricNumber field is 123456)
- If you just have these records in an array, or a linked list, you have to search them one by one
- On average, you have to look at half the entries in the database to find the one you want
- If you double the size of the database, you double the average time to look up an entry
- Terminology: this is called *linear time* or  $O(n)$  complexity



## A better approach

- Linear search is silly if you can place the records in order
- You're trying to find the telephone number of John Smith in the phone book
- Would you do this?
  - ① It isn't Aardvark, Aaron
  - ② It isn't Acton, Rebecca
  - ③ It isn't Anderson, Susan
  - ④ It isn't Atwater, James
  - ⑤ ... 8 million more unsuccessful searches because this is the Los Angeles directory
- That's silly!
- Open the book to the middle, notice that S is in the second half
- Open the book to the middle of the second half ...
- Each time you look at an entry in the book, you discard **half of the remaining possibilities**

# Binary search tree



- At every level: if a node contains  $x$ , then
  - ▶ every node in the left subtree is less than  $x$ , and
  - ▶ every node in the right subtree is greater than  $x$ .
- You can search the tree by starting at the root, and at every step you *know* whether to go left or right

# Algorithmic Complexity

- Complexity is concerned with **how the execution time grows as the size of the input grows**
- This is expressed as a function of the input size  $n$
- Normally we don't care about the *exact* function, and we use O-notation. Instead of a function like  $f(n) = 4.823 \times n$ , we just write  $f(n) = O(n)$ 
  - ▶  $O(1)$  — if input grows, the execution time remains unchanged. This is unrealistic: the program cannot even look at the input!
  - ▶  $O(n)$  — if the input is 5 times bigger, the execution time is 5 times bigger. This is the best you can hope for
  - ▶  $O(n^2)$  — if the input is 5 times bigger, the time is 25 times bigger

# Algorithm is more important than small optimisation

- Some programmers spend lots of effort trying to save one or two instructions in a piece of a program
  - ▶ But it doesn't matter much whether a program takes 2.00032 seconds or 2.00031 seconds
- It's much more important to use a suitable **algorithm**
  - ▶ On small data it doesn't make much difference
  - ▶ On large (realistic) data, a better algorithm makes a huge difference

# Complexity for search

- Ordered lists

- ▶ The Ordered Lists program has an operation to search a list for a value  $x$
- ▶ On average, you need to look through half of the data to find out whether  $x$  is present
- ▶ If the list were *not* ordered, you would need to look through *all* of the data to determine whether  $x$  is present
- ▶ So the ordered list makes the search about twice as fast
- ▶ But in either case, this is  $O(n)$  — if you double the data size, the average time is doubled

- Binary search tree

- ▶ The number of comparisons needed is roughly the height of the tree
- ▶ If the tree is *balanced*, the time complexity is  $O(\log n)$

# How much faster?

- With a linear data structure (array, linked list)
  - ▶ Each time you compare a database entry with your key, you eliminate one possibility
  - ▶ The time is proportional to the size of the database
  - ▶ It's called *linear time* — time =  $O(n)$
  - ▶ For 2 million records, you need a million comparisons
- With a binary search tree
  - ▶ Each time you compare a database entry with your key, you eliminate (on average) half of the possibilities
  - ▶ The time is proportional to *the logarithm of the size* of the database
  - ▶ It's called *log time* — time =  $O(\log n)$
  - ▶ For 2 million records, you need 21 comparisons
  - ▶ There's a saying: “logs come from trees”

# A common pitfall

- When you're writing a program, it's natural to test it with small data
- Even if the algorithm has bad complexity, the testing may be fast
- But then, when you run the program on real data, the execution time is intolerable
- That means going back and starting over again
- So it's a good idea to be aware of the complexity of your algorithm from the beginning

# How bad can complexity be?


Order of magnitude estimate of time for input of size  $n$

$n$	$\log n$	$n \log n$	$n^2$	$2^n$
1	1	1	1	2
10	3	30	100	1,000
100	7	700	10,000	1267650600228229401496703205376
1,000	10	10,000	1,000,000	> age of universe

- Lots of real problems have data size larger than 1,000
- Lots of algorithms have exponential complexity:  $2^n$



## tree



<https://xkcd.com/835/>