



University
of Glasgow | School of
Computing Science

Networks & Operating Systems Essentials

Dr Angelos Marnerides

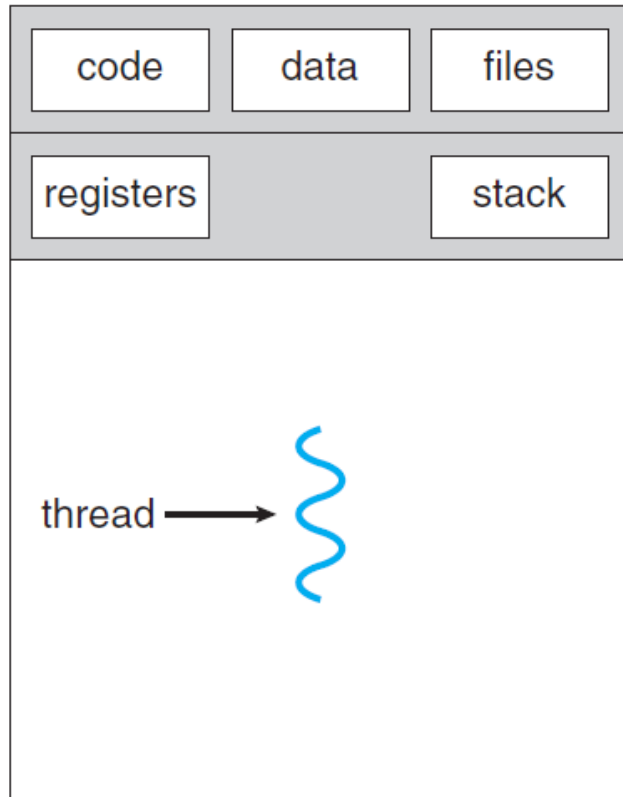
<angelos.marnerides@glasgow.ac.uk>

School of Computing Science

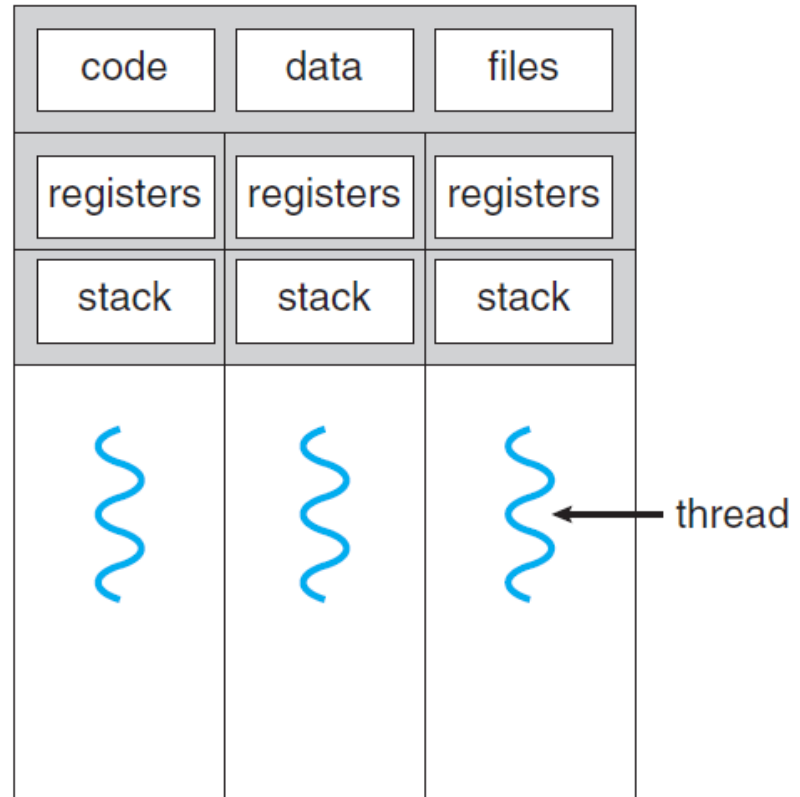
Concurrency vs Parallelism

- Parallelism: multiple tasks (threads/processes) execute at the same time
 - Requires multiple execution units (CPUs, CPU cores, CPU vCores, etc.)
- Concurrency: multiple tasks make progress over time
 - Time-sharing system

Threads



single-threaded process



multithreaded process

Source: A. Silberschatz, "Operating System Concepts", 9th Ed., 2012.

Threads vs Processes

- Threads share memory and resources by default
 - Processes require extra IPC mechanisms (shared memory, message passing) to be configured explicitly
- Threads are faster/more economical to create
 - Each process has its own copy of the in-memory data, hence process creation means page table duplication (possibly also data duplication, if copy-on-write not used)
- In older operating systems, it was much faster to context switch between threads than processes
 - No longer the case, as the kernel keeps the same info for both, hence context switching has almost identical cost
- Can have different schedulers for processes and threads
 - Depends on how threads are implemented (user threads vs kernel threads)
- If a single thread in a process crashes, the whole process crashes as well
 - If a process dies, other processes are unaffected

Aside: Amdahl's Law



- Assume you have a program, code such that:
 - Part of the code can only be executed serially (S)
 - Part of the code can be parallelised across N cores
- Then the maximum speed-up we can gain by going parallel is given by:

$$speedup = \frac{1}{\left(S + \frac{1-S}{N}\right)}$$

- Note: if N approaches infinity, then $speedup = 1/S$

A Banking Quiz...

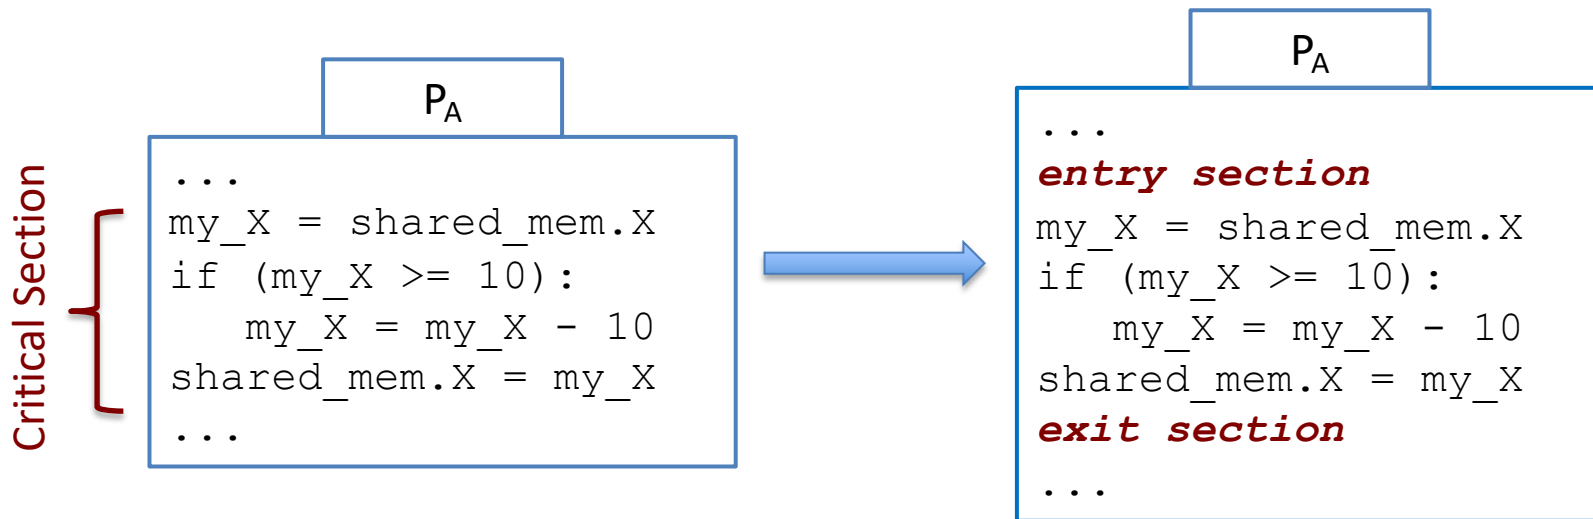
- Consider two processes P_A and P_B , communicating via shared memory
- Assume shared memory consists of a single integer $X = 20$



- What is the final value of X ?**
- P_A executes fully before P_B :
 - P_A 's Check for $X \geq 10$ succeeds $\rightarrow X = 10$; P_B 's: check for $X \geq 15$ fails $\rightarrow X = 10$
- P_B executes fully before P_A :
 - P_B 's check for $X \geq 15$ succeeds $\rightarrow X = 5$; P_A 's check for $X \geq 10$ fails $\rightarrow X = 5$
- But enter time-sharing/preemptive scheduling/parallel execution/...

A Banking Quiz...

- How would you alleviate this problem?
- Critical Section: A section of code where a process touches shared resources
 - Shared memory, common variables, shared database, shared file, ...
- Goal: No two processes in their critical section at the same time



The Critical Section problem

- Properties of acceptable solutions:
 1. **Mutual Exclusion**: No two processes in their critical section at the same time
 2. **Progress**: Entering one's critical section should only be decided by its contenders in due time (assuming no process already in its critical section); a process cannot immediately re-enter its critical section if other processes are waiting for their turn
 3. **Bounded waiting**: It should be impossible for a process to wait indefinitely to enter its critical section, if other processes are allowed to do so

Peterson's solution

P_A

```
...  
flag_A = True  
turn = B  
while flag_B == True and turn == B:  
    pass  
my_X = shared_mem.X  
if (my_X >= 10):  
    my_X = my_X - 10  
shared_mem.X = my_X  
flag_A = False  
...
```

P_B

```
...  
flag_B = True  
turn = A  
while flag_A == True and turn == A:  
    pass  
my_X = shared_mem.X  
if (my_X >= 10):  
    my_X = my_X - 10  
shared_mem.X = my_X  
flag_B = False  
...
```

- **Mutual Exclusion:** If P_A is in its critical section, then either flag_B = False (i.e., P_B is out of its critical section) or turn = A (i.e., P_B is waiting in the while loop)
- **Progress:** P_A cannot immediately reenter its critical section, as turn = B means P_B will be given the go next
- **Bounded waiting:** P_A will wait at most one turn before it can enter its critical section again

Peterson's solution

- Is that good enough?
 - **Busy waiting**: the waiting process eats up CPU cycles unnecessarily (a.k.a., spinlock)
 - **Memory reordering**: CPUs tend to reorder execution of mem accesses to avoid pipeline stalls
- Better solution: atomic instructions at the **hardware** level
 - “Test and set”
 - “Compare and swap”

value passed **by reference**

```
test_and_set(value):  
    my_value = value  
    value = True  
    return my_value
```

```
compare_and_swap(value, expected, new_value):  
    my_value = value  
    if value == expected:  
        value = new_value  
    return my_value
```

Peterson's solution revisited

test_and_set

lock {
...
while test_and_set(lock) == True:
 pass
my_X = shared_mem.X
if (my_X >= 10):
 my_X = my_X - 10
shared_mem.X = my_X
lock = False
...
unlock }

```
test_and_set(value):  
    my_value = value  
    value = True  
    return my_value
```

compare_and_swap

lock {
...
while compare_and_swap(lock, False, True) == True:
 pass
my_X = shared_mem.X
if (my_X >= 10):
 my_X = my_X - 10
shared_mem.X = my_X
lock = False
...
unlock }

```
compare_and_swap(value, expected, new_value):  
    my_value = value  
    if value == expected:  
        value = new_value  
    return my_value
```

- **Mutual Exclusion:** Only one process will execute `test_and_set/compare_and_swap` with the lock value originally being `False`
- **Progress/Bounded waiting:** There's nothing stopping a process from immediately re-entering its critical section!
- Extra: More elaborate/better solutions exist, but not covered here.

Beyond test_and_set/compare_and_swap

- Test_and_set/compare_and_swap work but are a bit clunky and a bit too low-level
- Enter *mutex locks* and *semaphores*
 - Internal state: a single integer value
 - For mutexes can only take values 0 or 1, for semaphores only values ≥ 0
 - API offers two *atomic* functions:

```
acquire(mutex):  
    while mutex == 0:  
        pass  
    mutex = 1
```

```
release(mutex):  
    mutex = 0
```

```
wait(semaphore):  
    while semaphore <= 0:  
        pass  
    semaphore -= 1
```

```
signal(semaphore):  
    semaphore += 1
```

- But still busy waiting/spinlocking?
- Can augment semaphores with a list of blocked processes each:
 - wait(semaphore) would instead add processes to said list if value is ≤ 0
 - signal(semaphore) would instead remove one process from said queue
- Common understanding today: use *mutex* for locking and *semaphores* for signaling

Discussion: atomic ops vs spinlocking

- Atomic operations provide a good solution to inter-process/thread synchronization
- But **how** does one do an atomic op in hardware?
 - **Disable interrupts** while atomic function is running
- How do you do that in a **multi-processing system**?
 - Good luck with that...
- **What then?**

Semaphores how to

- Set semaphore value to *number/size* of shared resources -- i.e., number of acceptable concurrent users of the resource
 - 1 if only 1 user is allowed, N for an N-sized queue, etc.
- Decrease (wait) the semaphore every time a resource is used
- Increase (signal) the semaphore every time a resource is released

Case study: bounded buffer problem

- Also known as the producer-consumer problem
- Assume a list where items are placed by producers, and removed by consumers
- Assume we want our list to never contain more than N items
- Goal: Allow producers of items to add them to the list, but have them wait first if the list is full
- Goal: Allow consumers of items to remove an item from the list, but have them wait first if the list is empty

```
semaphore mutex = 1  
semaphore empty = N  
semaphore full = 0
```

```
producer():  
    while True:  
        # Produce an item  
        wait(empty)  
        wait(mutex)  
        # Add item to list  
        signal(mutex)  
        signal(full)
```

```
consumer():  
    while True:  
        wait(full)  
        wait(mutex)  
        # Remove an item from list  
        signal(mutex)  
        signal(empty)  
        # Do stuff with item
```

Case study: readers-writers problem

- Assume a variable shared among many processes, some of which only read its value (readers) while others also need to update it (writers)
- Goal: When a writer accesses the variable, no other process should be able to either read or update it
- Goal: When a reader accesses the variable, more readers can also access it, but writers should wait until no reader accesses it

```
semaphore rw_mutex = 1
semaphore mutex    = 1
int read_count     = 0
```

```
writer():
    while True:
        wait(rw_mutex)
        # Update the value
        signal(rw_mutex)
```

```
reader():
    while True:
        wait(mutex)
        read_count += 1
        if (read_count == 1)
            wait(rw_mutex)
        signal(mutex)
        # Read value
        wait(mutex)
        read_count -= 1
        if (read_count == 0)
            signal(rw_mutex)
        signal(mutex)
```


Matters of life and death...

- What would happen with writers if readers keep on arriving at the system?
 - Writers would “starve”
- **Starvation**: Processes/threads unable to enter their critical section because of “greedy” contenders
 - Think: trying to get on an extremely busy motorway with no one giving you some space
- **Livelock**: special case of starvation where competing parties both try to “avoid” each other at the same time
 - Think: bumping into a person in a corridor, then both going left/right at the same time only to bump into each other again
- **Priority inversion**: special case where lower priority processes can keep higher priority processes waiting
 - Think: having to sleep but being kept awake by social media notifications...

Matters of life and death...

- What would happen in the producer/consumer problem if order of locking/unlocking mutex and empty/full was reversed?

```
semaphore mutex = 1  
semaphore empty = N  
semaphore full = 0
```

```
producer():  
    while True:  
        # Produce an item  
        wait(mutex)  
        wait(empty)  
        # Add item to list  
        signal(full)  
        signal(mutex)
```

```
consumer():  
    while True:  
        wait(mutex)  
        wait(full)  
        # Remove an item from list  
        signal(empty)  
        signal(mutex)  
        # Do stuff with item
```

- Assume a consumer goes first...
 - `mutex` → locked; consumer waiting on `full`
 - producer waiting on `mutex`
- **Deadlock**: all parties of a group waiting indefinitely for another party (incl. themselves) to take action

Beyond semaphores...

- Semaphores/mutexes allow for mutual exclusion, but once a process is blocked that's it
- Enter **Monitors**
 - Combination of semaphores and **condition variables**
 - Each condition variable “associated” with a semaphore
 - Allows for processes to have both mutual exclusion, and wait (block) on a condition
 - `cond_wait(condvar, mutex)`: unlock the mutex to wait for a condition, then atomically reacquire the mutex when condition is met
 - `cond_signal(condvar)`: unblock one of the processes waiting on condition


```
semaphore mutex = 1
condvar   empty = N
condvar   full  = 0
list      items = 0
```

```
producer():
    while True:
        # Produce an item
        wait(mutex)
        while len(items) == N:
            cond_wait(empty, mutex)
        # Add item to list
        cond_signal(full)
        signal(mutex)
```

```
consumer():
    while True:
        wait(mutex)
        while len(items) == 0:
            cond_wait(full, mutex)
        # Remove an item from list
        cond_signal(empty)
        signal(mutex)
        # Do stuff with item
```

Food for thought: dining philosophers

- Philosophers sitting on a round table, plate in front of them, but only as many chopsticks as there are philosophers (see image)
- Philosophers spend most time thinking, but every now and then get hungry...
- One needs two chopsticks to eat and won't grab a chopstick from their neighbour's hand
- Philosophers can only pick up one chopstick at a time
- When done eating, they'll put the chopsticks back on the table



Source: A. Silberschatz, "Operating System Concepts", 9th Ed., 2012.

Recommended Reading

- Silberschatz et. al., "Operating Systems Essentials", Chapter 6, sections 6.3,6.4, 6.5