

Algorithms and Data Structures 2

15 - Red-black trees

Dr Michele Sevegnani

School of Computing Science
University of Glasgow


michele.sevegnani@glasgow.ac.uk

Outline

- **Recap**
- **Red-black trees**
 - Definition
 - Representation
 - Properties
 - Insertion
- **Comparison with AVL trees**

Recap

- Each of the basic operations on a binary search tree runs in $O(h)$ time
 - h is the height of the tree
- However, the height varies as items are inserted and deleted
- Try to insert elements **1,2,3,4,5** (in this order) into an empty BST
 - Unbalanced tree with height **4**
 - Height is $O(n)$ in unbalanced trees



Self-balancing trees


- **Several extensions to the basic BST definition have been introduced to keep the height small as items are dynamically inserted and deleted**
 - Red-black trees (in this lecture)
 - AVL trees
 - B-trees
- **A common method to keep the tree balanced is to perform rotations after each deletion and insertion**
 - Local operation in a search tree that preserves the binary-search-tree property

Red-Black trees

- **Self-balancing binary search trees introduced by Bayer in 1972**
 - Initially known as **symmetric binary B-trees**
 - Bayer, Rudolf. "Symmetric binary B-trees: Data structure and maintenance algorithms." Acta informatica 290-306
- **Red/black convention introduced by Sedgewick in 1978**
 - Guibas, Leo J., and Robert Sedgewick. "A dichromatic framework for balanced trees." 19th Annual Symposium on Foundations of Computer Science. IEEE
- **Basic operations on a red-black tree take $O(\log n)$ time in the worst case**

Definition

- A red-black tree is a binary search tree with an extra attribute **colour**, which can be either **RED** or **BLACK** and satisfies the **red-black properties**



- **All leaves (external nodes) of the tree contain the value NIL**
- **Internal nodes are the normal nodes containing regular keys**

Example


- In this representation
 - Pointer attribute p is not shown

- Each **NIL** leaf is **black**

- Each internal node x is marked with its **black-height** ($bh(x)$)

- black
- x
-

ADS 2, 2021



Sentinel

- For a tree **T**, **NIL** nodes are represented by a single **sentinel T.nil**

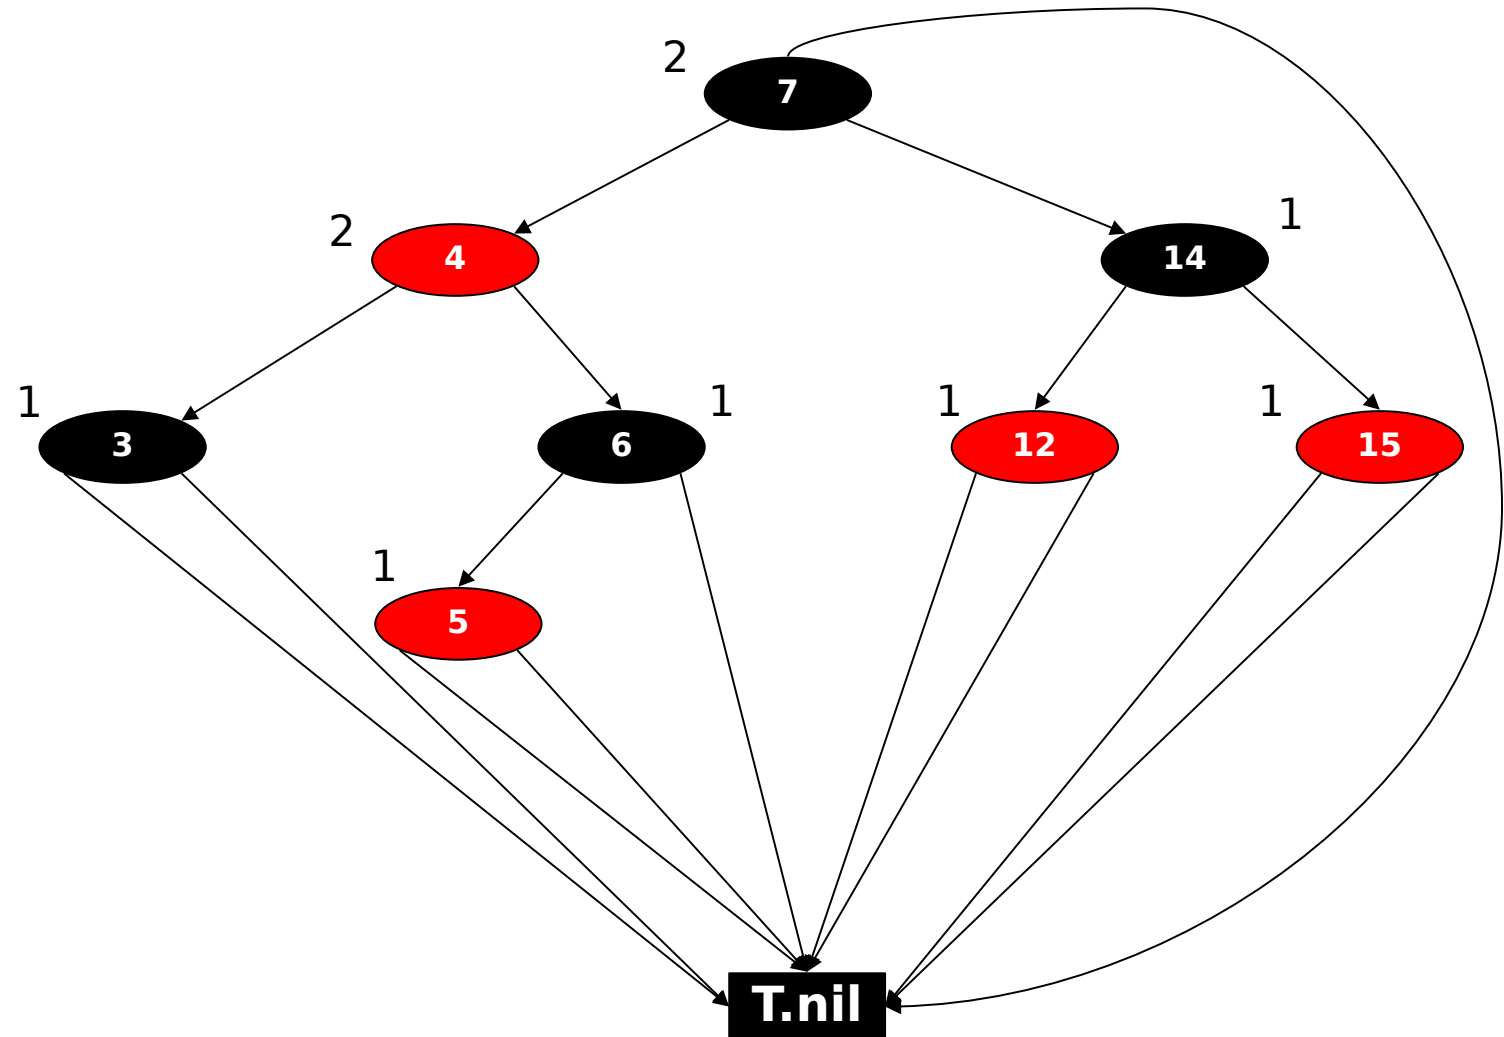
- Colour attribute is **BLACK**

- **NIL**

-


-

- The root's parent is also the sentinel



Standard representation

- Sentinel is omitted and only internal nodes are shown
- Black-heights are omitted



Properties

- A red-black tree with n internal nodes has height h at most $2 \log (n + 1)$
- Proof is in two steps
 1. Prove that the subtree rooted at x contains at least $2^{bh(x)} - 1$ internal nodes
 2. Prove the bound of the height: $h \leq 2 \log(n + 1)$

Proof of 1

1. The subtree rooted at x contains at least $2^{bh(x)} - 1$ internal nodes

- By induction on $h(x)$, the height of x
- **Base $h(x) = 0$**
 - Then x is the leaf $T.nil$ and its subtree contains $2^{bh(x)} - 1 = 2^0 - 1 = 0$
- **Inductive step**
 - Consider an internal node x with positive height and two children
 - Each child has black-height $bh(x)$ (if red) or $bh(x) - 1$ (if black)
 - $2^{bh(x)-1} - 1$
 - $x \quad (2^{bh(x)-1} - 1) + (2^{bh(x)-1} - 1) + 1 = 2^{bh(x)} - 1$

Proof of 2

2. Prove $h \leq 2 \log(n + 1)$

- By property 4, at least half of the nodes on a any simple path from the root (excluded) to the leaf, must be **black**

$$- \quad \text{bh}(x) \quad - \quad h/2. \quad n \geq 2^{\text{bh}(x)} - 1$$

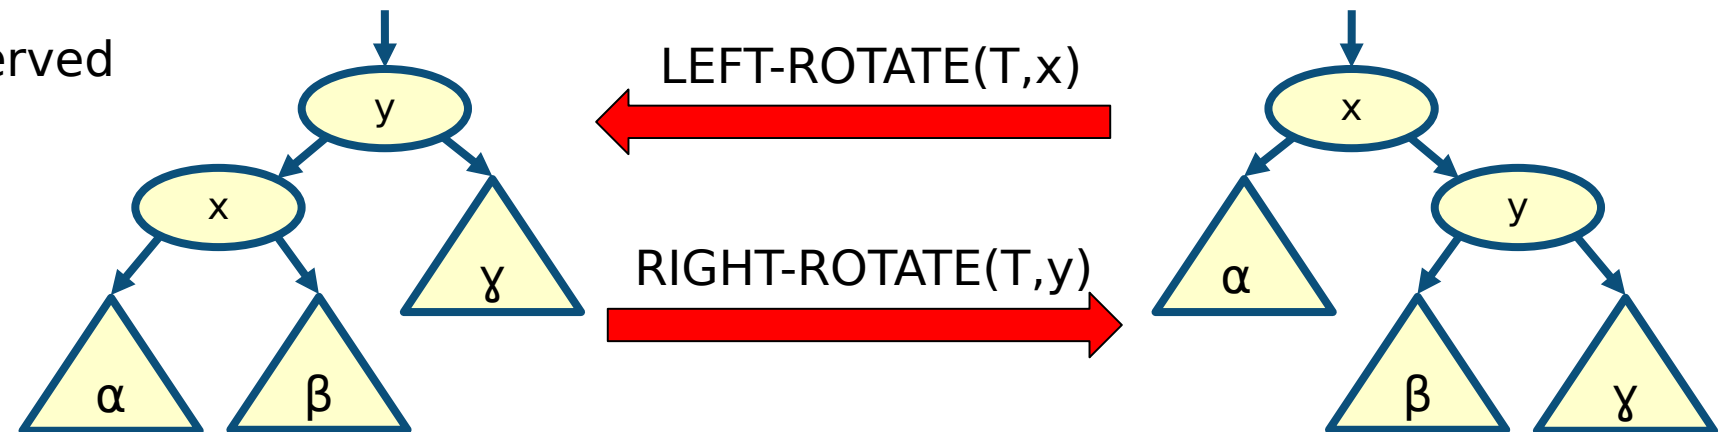


$$h \leq 2 \log (n + 1)$$

- This guarantees that operations on a red-black tree take $O(\log n)$ time in the worst case

Rotations

- **Insertions and deletions may violate the red-black properties**
- **To restore these properties, we must**
 1. Change the **colours** of some of the nodes (later in this lecture)
 2. Change the pointer structure through **left** and **right rotations**
- **Example rotations**
 - α , β , and γ are arbitrary subtrees
 - BST property is preserved





Left rotation

- The left rotation “pivots” around the link from **x** to **y**
- **Assumptions**
 - Right child of **x** is not **T.nil**
 - The root’s parent is **T.nil**
- **Running time is $O(1)$**
- The pseudocode for **RIGHT-ROTATE** is symmetric


```
LEFT-ROTATE(T, x)
  y := x.right
  x.right := y.left
  if y.left != NIL
    y.left.p := x
  y.p := x.p
  if x.p = T.nil
    T.root := y
  elseif x = x.p.left
    x.p.left := y
  else x.p.right := y
  y.left := x
  x.p := y
```

Example

- Try to left rotate on node with key **3**
 - $\alpha = \beta = T.nil$
 - γ is the subtree rooted at 6





LEFT-ROTATE(T,x)




Example

- **Left rotation on node with key 3**

- $\alpha = \beta = T.nil$
- γ is the subtree rooted at 6



LEFT-ROTATE(T,x)



Insertion

- Like insertion for BST with a **few differences**
 1. Sentinel **T.nil** replaces **NIL**
 2. New node **z** is coloured **red**
 3. **FIXUP** is called at the end to restore the red-black properties
- The definition of **FIXUP** is quite involved
 - We first analyse the violations to the red-black properties introduced by inserting a new **red** node

```
INSERT(T, z)
  y := T.nil
  x := T.root
  while x != NIL
    y := x
    if z.key < x.key
      x := x.left
    else x := x.right
  z.p := y
  if y = T.nil
    T.root := z
  elseif z.key < y.key
    y.left := z
  else y.right := z
  z.left := T.nil
  z.right := T.nil
  z.colour := RED
  FIXUP(T, z)
```

Red-black properties violations

- **When a new red node *z* is inserted**
 - Properties 1,3 and 5 still hold
 - Property 2 is violated if *z* is the root
 - Property 4 is violated if *z*'s parent is red


Red-black properties

1. Every node is either red or black




Example violation

- Try to insert **6** in the red-black tree below




Red-black properties

1. Every node is either **red** or **black**



Example violation

- Violation of Property 4



Red-black properties

1. Every node is either **red** or **black**




FIXUP

- **Operation that restores the red-black tree property after an insertion or a deletion**
 - **Recolour** nodes and perform **rotations**
- **It consists of a while loop and three different cases**
- **We cover them one by one**

```
FIXUP(T,z)
  while z.p.colour = RED
    if z.p = z.p.p.left
      y := z.p.p.right
      if y.colour = RED
        z.p.colour := BLACK // case 1
        y.colour := BLACK // case 1
        z.p.p.colour := RED // case 1
        z := z.p.p // case 1
      else
        if z = z.p.right
          z := z.p // case 2
          LEFT-ROTATE(T,z) // case 2
          z.p.colour := BLACK // case 3
          z.p.p.colour := RED // case 3
          RIGHT-ROTATE(T,z.p.p) // case 3
        else
          ... // symmetric
  T.root.colour := BLACK
```


Case 1: z's uncle y is red

- **z** is a right child (same when **z** is a left child)
- Each of the subtrees **α** , **β** , **γ** , **δ** , and **ϵ** has a **black root**, and each has the **same black-height**
- The while loop continues with node **z's grandparent z.p.p** as the new **z**




Case 2: z's uncle y is black and z is a right child

- Transform into case 3 with a single left rotation on **z's** parent



Case 3: z's uncle y is black and z is a left child

- Perform two colour changes (**z.p** and **z.p.p**) and a right rotation on **z.p.p**
- The while loop then terminates, because there are no longer two **red** nodes in a row



Example

- Insert keys **12,7,3,15,4,6,14,5** in the empty red-black tree **T**

Example


- Insert keys **12,7,3,15,4,6,14,5** in the empty red-black tree **T**
 - Insert 12
 - Loop not executed

12

```
FIXUP(T,z)
  while z.p.colour = RED
    if z.p = z.p.p.left
      y := z.p.p.right
      if y.colour = RED
        z.p.colour := BLACK
        y.colour := BLACK
        z.p.p.colour := RED
        z := z.p.p
      else
        if z = z.p.right
          z := z.p
          LEFT-ROTATE(T,z)
          z.p.colour := BLACK
          z.p.p.colour := RED
          RIGHT-ROTATE(T,z.p.p)
        else
          ...
  T.root.colour := BLACK
```

Example


- Insert keys **12,7,3,15,4,6,14,5** in the empty red-black tree **T**
 - Insert 7
 - Loop not executed



```
FIXUP(T,z)
  while z.p.colour = RED
    if z.p = z.p.p.left
      y := z.p.p.right
      if y.colour = RED
        z.p.colour := BLACK
        y.colour := BLACK
        z.p.p.colour := RED
        z := z.p.p
      else
        if z = z.p.right
          z := z.p
          LEFT-ROTATE(T,z)
          z.p.colour := BLACK
          z.p.p.colour := RED
          RIGHT-ROTATE(T,z.p.p)
        else
          ...
  T.root.colour := BLACK
```

Example


- Insert keys **12,7,3,15,4,6,14,5** in the empty red-black tree **T**
 - Insert 3
 - Case 3: RIGHT-ROTATE



```
FIXUP(T,z)
  while z.p.colour = RED
    if z.p = z.p.p.left
      y := z.p.p.right
      if y.colour = RED
        z.p.colour := BLACK
        y.colour := BLACK
        z.p.p.colour := RED
        z := z.p.p
      else
        if z = z.p.right
          z := z.p
          LEFT-ROTATE(T,z)
          z.p.colour := BLACK
          z.p.p.colour := RED
          RIGHT-ROTATE(T,z.p.p)
        else
          ...
  T.root.colour := BLACK
```

Example


- Insert keys **12,7,3,15,4,6,14,5** in the empty red-black tree **T**
 - After FIXUP



```
FIXUP(T,z)
  while z.p.colour = RED
    if z.p = z.p.p.left
      y := z.p.p.right
      if y.colour = RED
        z.p.colour := BLACK
        y.colour := BLACK
        z.p.p.colour := RED
        z := z.p.p
      else
        if z = z.p.right
          z := z.p
          LEFT-ROTATE(T,z)
          z.p.colour := BLACK
          z.p.p.colour := RED
          RIGHT-ROTATE(T,z.p.p)
        else
          ...
  T.root.colour := BLACK
```

Example


- Insert keys **12,7,3,15,4,6,14,5** in the empty red-black tree **T**
 - Insert 15
 - Case 1: push blackness down from grandparent



```
FIXUP(T,z)
while z.p.colour = RED
  if z.p = z.p.p.left
    y := z.p.p.right
    if y.colour = RED
      z.p.colour := BLACK
      y.colour := BLACK
      z.p.p.colour := RED
      z := z.p.p
    else
      if z = z.p.right
        z := z.p
        LEFT-ROTATE(T,z)
      z.p.colour := BLACK
      z.p.p.colour := RED
      RIGHT-ROTATE(T,z.p.p)
  else
    ...
T.root.colour := BLACK
```

Example


- Insert keys **12,7,3,15,4,6,14,5** in the empty red-black tree **T**
 - Insert 15
 - Case 1: push blackness down from grandparent



```
FIXUP(T,z)
while z.p.colour = RED
  if z.p = z.p.p.left
    y := z.p.p.right
    if y.colour = RED
      z.p.colour := BLACK
      y.colour := BLACK
      z.p.p.colour := RED
      z := z.p.p
    else
      if z = z.p.right
        z := z.p
        LEFT-ROTATE(T,z)
      z.p.colour := BLACK
      z.p.p.colour := RED
      RIGHT-ROTATE(T,z.p.p)
  else
    ...
T.root.colour := BLACK
```

Example


- Insert keys **12,7,3,15,4,6,14,5** in the empty red-black tree **T**
 - Restore root



```
FIXUP(T,z)
while z.p.colour = RED
    if z.p = z.p.p.left
        y := z.p.p.right
        if y.colour = RED
            z.p.colour := BLACK
            y.colour := BLACK
            z.p.p.colour := RED
            z := z.p.p
        else
            if z = z.p.right
                z := z.p
                LEFT-ROTATE(T,z)
            z.p.colour := BLACK
            z.p.p.colour := RED
            RIGHT-ROTATE(T,z.p.p)
        else
            ...
T.root.colour := BLACK
```


Example


- Insert keys **12,7,3,15,4,6,14,5** in the empty red-black tree **T**
 - Insert 4
 - No violation



```
FIXUP(T,z)
while z.p.colour = RED
  if z.p = z.p.p.left
    y := z.p.p.right
    if y.colour = RED
      z.p.colour := BLACK
      y.colour := BLACK
      z.p.p.colour := RED
      z := z.p.p
    else
      if z = z.p.right
        z := z.p
        LEFT-ROTATE(T,z)
      z.p.colour := BLACK
      z.p.p.colour := RED
      RIGHT-ROTATE(T,z.p.p)
  else
    ...
T.root.colour := BLACK
```

Example


- Insert keys **12,7,3,15,4,6,14,5** in the empty red-black tree **T**
 - Insert 6
 - Case 3 (symmetric): LEFT-ROTATE



```
FIXUP(T,z)
while z.p.colour = RED
  if z.p = z.p.p.left
    y := z.p.p.right
    if y.colour = RED
      z.p.colour := BLACK
      y.colour := BLACK
      z.p.p.colour := RED
      z := z.p.p
    else
      if z = z.p.right
        z := z.p
        LEFT-ROTATE(T,z)
      z.p.colour := BLACK
      z.p.p.colour := RED
      RIGHT-ROTATE(T,z.p.p)
  else
    ...
T.root.colour := BLACK
```

Example


- Insert keys **12,7,3,15,4,6,14,5** in the empty red-black tree **T**
 - After FIXUP



```
FIXUP(T,z)
while z.p.colour = RED
  if z.p = z.p.p.left
    y := z.p.p.right
    if y.colour = RED
      z.p.colour := BLACK
      y.colour := BLACK
      z.p.p.colour := RED
      z := z.p.p
    else
      if z = z.p.right
        z := z.p
        LEFT-ROTATE(T,z)
      z.p.colour := BLACK
      z.p.p.colour := RED
      RIGHT-ROTATE(T,z.p.p)
  else
    ...
T.root.colour := BLACK
```

Example


- Insert keys **12,7,3,15,4,6,14,5** in the empty red-black tree **T**
 - Insert 14
 - Case 2 (symmetric): RIGHT-ROTATE



```
FIXUP(T,z)
while z.p.colour = RED
  if z.p = z.p.p.left
    y := z.p.p.right
    if y.colour = RED
      z.p.colour := BLACK
      y.colour := BLACK
      z.p.p.colour := RED
      z := z.p.p
    else
      if z = z.p.right
        z := z.p
        LEFT-ROTATE(T,z)
        z.p.colour := BLACK
        z.p.p.colour := RED
        RIGHT-ROTATE(T,z.p.p)
      else
        ...
T.root.colour := BLACK
```

Example


- Insert keys **12,7,3,15,4,6,14,5** in the empty red-black tree **T**
 - Case 3 (symmetric): LEFT-ROTATE



```
FIXUP(T,z)
while z.p.colour = RED
  if z.p = z.p.p.left
    y := z.p.p.right
    if y.colour = RED
      z.p.colour := BLACK
      y.colour := BLACK
      z.p.p.colour := RED
      z := z.p.p
    else
      if z = z.p.right
        z := z.p
        LEFT-ROTATE(T,z)
        z.p.colour := BLACK
        z.p.p.colour := RED
        RIGHT-ROTATE(T,z.p.p)
      else
        ...
T.root.colour := BLACK
```

Example


- Insert keys **12,7,3,15,4,6,14,5** in the empty red-black tree **T**
 - After FIXUP



```
FIXUP(T,z)
while z.p.colour = RED
  if z.p = z.p.p.left
    y := z.p.p.right
    if y.colour = RED
      z.p.colour := BLACK
      y.colour := BLACK
      z.p.p.colour := RED
      z := z.p.p
    else
      if z = z.p.right
        z := z.p
        LEFT-ROTATE(T,z)
        z.p.colour := BLACK
        z.p.p.colour := RED
        RIGHT-ROTATE(T,z.p.p)
      else
        ...
T.root.colour := BLACK
```

Example


- Insert keys **12,7,3,15,4,6,14,5** in the empty red-black tree **T**
 - Insert 5
 - Case 1: push blackness down from grandparent



```
FIXUP(T,z)
while z.p.colour = RED
  if z.p = z.p.p.left
    y := z.p.p.right
    if y.colour = RED
      z.p.colour := BLACK
      y.colour := BLACK
      z.p.p.colour := RED
      z := z.p.p
    else
      if z = z.p.right
        z := z.p
        LEFT-ROTATE(T,z)
      z.p.colour := BLACK
      z.p.p.colour := RED
      RIGHT-ROTATE(T,z.p.p)
  else
    ...
T.root.colour := BLACK
```

Example

- Insert keys **12,7,3,15,4,6,14,5** in the empty red-black tree **T**
 - After FIXUP



```
FIXUP(T,z)
while z.p.colour = RED
  if z.p = z.p.p.left
    y := z.p.p.right
    if y.colour = RED
      z.p.colour := BLACK
      y.colour := BLACK
      z.p.p.colour := RED
      z := z.p.p
    else
      if z = z.p.right
        z := z.p
        LEFT-ROTATE(T,z)
      z.p.colour := BLACK
      z.p.p.colour := RED
      RIGHT-ROTATE(T,z.p.p)
  else
    ...
T.root.colour := BLACK
```


Analysis of INSERT

- Since the height of a red-black tree on n nodes is $O(\log n)$ the first part of **INSERT** takes $O(\log n)$ time
- In **FIXUP**, the while loop repeats only if **case 1** occurs, and then the pointer **z** moves two levels up the tree
 - The total number of times the while loop can be executed is $O(\log n)$
- **Therefore, INSERT takes a total of $O(\log n)$ time**
 - Observe that it never performs more than two rotations, since the while loop terminates if case 2 or case 3 is executed

Deletion of node x

- **Complex operation consisting of four cases**

1. x's sibling w is red

2. x's sibling w is black w black

■ x w black w red w black

■ x w black w red


- **Needed to fix violations of property 5**
- **It takes $O(\log n)$ time and performs at most three rotations**
- **Not part of the course**

AVL tree

- **G. Adelson-Velskii and E.M. Landis, "An algorithm for the organization of information." Doklady Akademii Nauk SSSR, 146:263-266, 1962**
- **An AVL tree is a binary search tree that is height balanced: for each node x , the heights of the left and right subtrees of x differ by at most 1**
- **To implement an AVL tree, we maintain an extra attribute in each node: $x.h$ is the height of node x**
- **$T.root$ points to the root node**

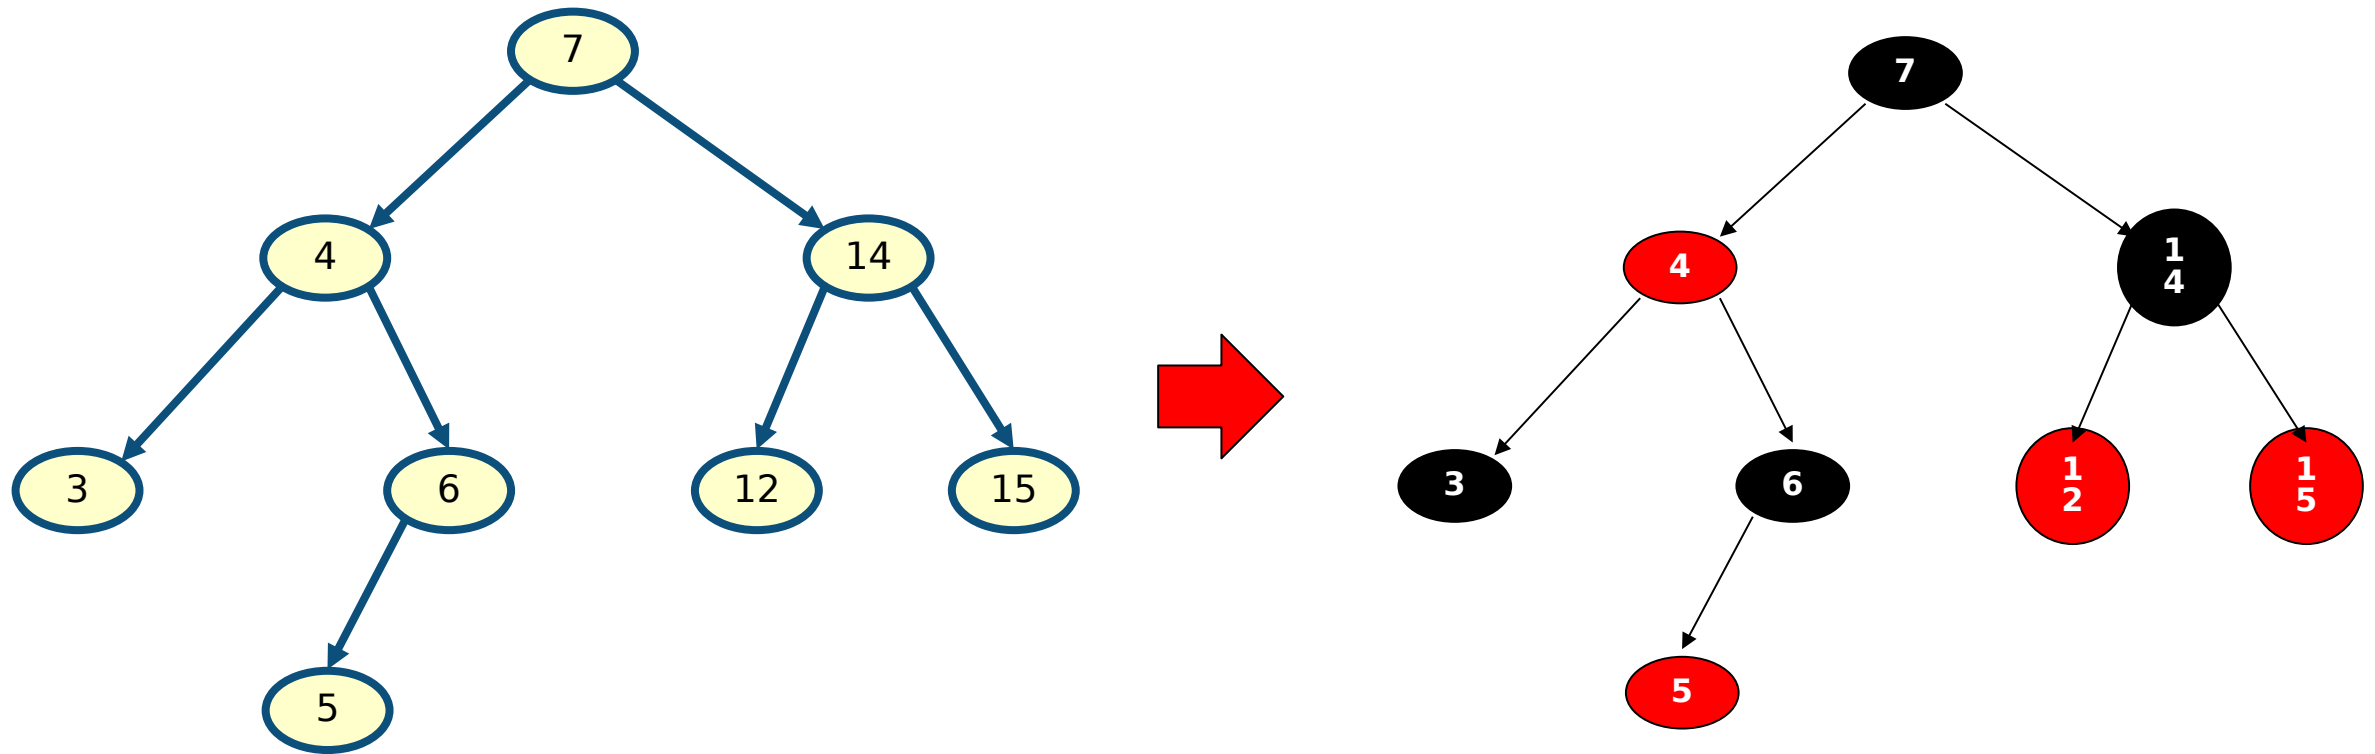
Example

- Height at each node is marked




AVL trees are a subset of red-black trees

- AVL trees can be coloured red-black
- There are red-black trees which are not AVL balanced



Example red-black tree, but not AVL tree

- Height of left and right subtrees of **7** differ by more than **1**
- **AVL** requires a right rotation on **7**



Comparison

- **AVL trees are more rigidly balanced than red-black trees**
 - $h \leq 2 \log (n + 1)$ in red-black trees
 - $h \leq 1.44 \log n$ in AVL trees
- **AVL trees provide faster lookups than red-black tree**
 - Used in databases
- **Red-black trees provide faster insertion and removal operations than AVL trees as fewer rotations are performed**
 - Used in real-time applications and the current Linux kernel scheduler
 - Used in most libraries to implement common ADTs such as Set, Multiset, and Map

Summary

- **Red-black trees**

- Definition
- Representation
- Properties
- Insertion

- **Comparison with AVL trees**