

Algorithms and Data Structures 2

16 - B-trees

Dr Michele Sevegnani

School of Computing Science
University of Glasgow

michele.sevegnani@glasgow.ac.uk

Outline

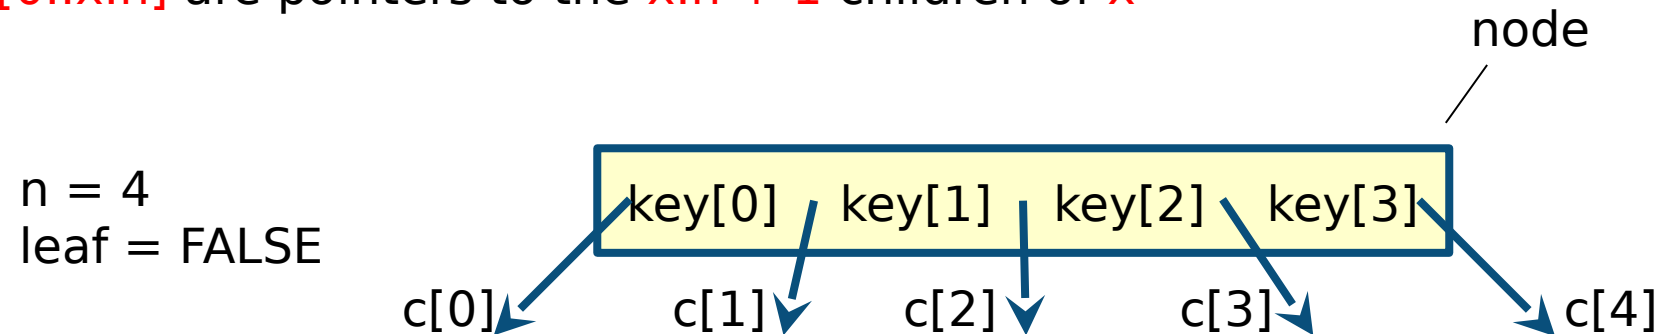
- **Definition**
- **Motivation**
- **Properties**
- **Operations**
 - Search
 - Insertions
- **Variants**
- **Applications**

B-trees

- **Balanced search trees introduced by Bayer in 1972**
 - Same inventor of red-black trees!
 - Bayer, Rudolf, and Edward M. McCreight. "Organization and Maintenance of Large Ordered Indices." Acta Informatica 1 (1972): 173-189.
- **Designed for big data sets stored on disks or other secondary storage devices**
 - Number of I/O operations is minimised
- **B-tree nodes may have more than two children**
 - Typically thousands depending on the physical characteristic of the hard disk
 - Height is $O(\log n)$ but usually much less than that of a red-black tree (example later in this lecture)

Definition

- **A node x in a B-tree has the following attributes**
 - $x.n$ is the number of keys stored in the node
 - $x.key[0..x.n-1]$ are the $x.n$ keys stored in nondecreasing order
 - $x.leaf$ is **TRUE** if x is a leaf, **FALSE** otherwise
 - $x.c[0..x.n]$ are pointers to the $x.n + 1$ children of x




- **An attribute $T.root$ points to the root of B-tree T**
- **We call a fixed integer $t \geq 2$ the minimum degree (number of children) of a B-tree**

Definition (cont.)

- **A B-tree is a rooted tree satisfying the following properties**


1. Leaves have no children (**x.c** is **NIL**)
2. The keys separate the ranges of keys stored in each subtree (see example below)



3. All leaves have the same depth (which is the tree height)
 4. Every node other than the root must have at least **t - 1** keys
 5. Every node may contain at most **2t - 1** keys
- } Bounds on the number of keys

Example

- **Minimum degree is $t = 2$**
 - Sometimes called **2-3-4 trees**, as every internal node has either **2**, **3**, or **4** children
- **The tree stores **18** keys and has height **2****
 - A red-black tree of height **5** is required to store **18** keys



Worst-case height of a B-tree

- If $n \geq 1$, then for any n -key B-tree T of height h and minimum degree $t \geq 2$

- **Proof**

- The root of T contains at least **one** key, while all other nodes have at least $t - 1$ keys
- Therefore, T has at least **2** nodes at depth **1**, $2t$ nodes at depth **2**, $2t^2$ nodes at depth **3**, ...
- In general, we have $2t^{h-1}$ nodes at depth h
- The total number of **nodes** is obtained by summing the nodes at each depth
- Therefore, the number of **keys** n satisfies the following inequality

- Take the **base- t log** on both sides to conclude the proof


Geometric series

Motivation for B-trees

- **When a data structure is stored in secondary memory accessing a single node may take up to 8-10 milliseconds**
 - More than 100k memory accesses can be performed in the same time span
 - Primary memory access time is 50 nanoseconds
- **Each access retrieves a page of bits (2^{11} - 2^{14} bytes) not just one item to amortise the time spent waiting for mechanical movements**
- **To minimise the number of disk access, B-trees minimise the number of nodes**
 - Height is reduced by allowing high branching factors
 - Bigger nodes, slower to process, but faster overall as disk access is very slow

Example

- A B-tree with branching factor **1001** and height **2** can store over **one billion** keys



1 node 1000 keys

1001 nodes $1001 * 10^3$ keys

1001^2 nodes $1001^2 * 10^3$ keys

- Keeping the root node in memory allows us to access any key by making at most 2 disk accesses

ADS 2, 2021 To minimise I/O, a B-tree node is usually as large as a **whole** disk page

– Typical branching factors range from **50** to **2000**

Modelling I/O

- **Typically the amount of data to be stored in a B-tree is so large that all the data cannot fit into main memory at once**
 - The B-tree algorithms copy selected **pages** from disk to main memory as needed and write back onto disk the pages that have changed
 - We assume the operating system **flushes** from main memory pages no longer in use
- **We explicitly model disk operations in the pseudocode**
 - **DISK-READ(x)** reads object x into main memory
 - **DISK-WRITE(x)** save object x onto disk
 - **ALLOCATE-NODE()** allocates one disk page in $O(1)$ time

Operations

- **The root of the B-tree is always stored in the main memory**
 - No **DISK-READ** on the root needed
 - A **DISK-WRITE** of the root is needed whenever the root is modified
- **Any nodes that are passed as parameters must already have had a **DISK-READ** operation performed on them**
- **We study the **one-pass** version of the B-tree operations**
 - Proceed downward from the root of the tree, without having to back up
 - **SEARCH**
 - **CREATE**
 - **INSERT**
 - DELETE is not part of the course

Search

- **Similar to recursive BST search**
 - Instead of doing a **binary decision** at each node (left/right), we do a **multiway decision** according to the number of children of the node
- **Top level call is SEARCH(T.root,k)**
- **Return a node x and an index i such that $x.key[i] = k$**
- **Return NIL if k not found**

```
SEARCH(x,k)
  i := 0
  while i < x.n and k > x.key[i]
    i := i + 1
  if i < x.n and k = x.key[i]
    return (x,i)
  elseif x.leaf
    return NIL
  else
    DISK-READ(x,c[i])
    return SEARCH(x.c[i],k)
```

Analysis


- **n** is the number of keys and **t** is the minimum degree
- **SEARCH** performs **$O(\log_t n)$** disk accesses
- **Total CPU time is $O(t \log_t n)$**
 - While loop takes **$O(t)$** time since $x.n < 2t$

```
SEARCH(x,k)
  i := 0
  while i < x.n and k > x.key[i]
    i := i + 1
  if i < x.n and k = x.key[i]
    return (x,i)
  elseif x.leaf
    return NIL
  else
    DISK-READ(x,c[i])
    return SEARCH(x.c[i],k)
```

Example

- Find key 16

- $x = T.\text{root}$




```
SEARCH(x,k)
  i := 0
  while i < x.n and k > x.key[i]
    i := i + 1
  if i < x.n and k = x.key[i]
    return (x,i)
  elseif x.leaf
    return NIL
  else
    DISK-READ(x,c[i])
    return SEARCH(x.c[i],k)
```

Example

- Find key 16

- 16 > 5 and 16 > 11
- $i = 2$
- Retrieve next node from disk
- Recourse



SEARCH(x,k)

```
i := 0
while i < x.n and k > x.key[i]
    i := i + 1
if i < x.n and k = x.key[i]
    return (x,i)
elseif x.leaf
    return NIL
else
    DISK-READ(x,c[i])
    return SEARCH(x.c[i],k)
```


Example

- **Find key 16**

- $15 < 16 < 21$
- $i = 1$
- Retrieve next node from disk
- Recourse

SEARCH(x,k)


```
i := 0
while i < x.n and k > x.key[i]
    i := i + 1
if i < x.n and k = x.key[i]
    return (x,i)
elseif x.leaf
    return NIL
else
    DISK-READ(x,c[i])
    return SEARCH(x.c[i],k)
```



Example

- **Find key 16**
 - $16 = x.\text{key}[0]$
 - Return a pointer to **x** (the green node) and index **0**

```
SEARCH(x,k)
  i := 0
  while i < x.n and k > x.key[i]
    i := i + 1
  if i < x.n and k = x.key[i]
    return (x,i)
  elseif x.leaf
    return NIL
  else
    DISK-READ(x,c[i])
    return SEARCH(x.c[i],k)
```



Creating an empty B-tree

- Requires **$O(1)$** disk operations
- **$O(1)$** CPU time

```
CREATE(T)
  x := ALLOCATE-NODE()
  x.leaf := TRUE
  x.n := 0
  DISK-WRITE(x)
  T.root := x
```

Insertion

- As with BSTs, we search for the **leaf** position at which to insert a new key
- In order not to violate the B-tree properties, a new key is added into an **existing leaf** node
- If a leaf node is full, we call **SPLIT-CHILD** to split the leaf into two nodes and move the median key up the leaf's parent
- If the parent is also full, we must split it before we can insert the new key, and thus we could end up splitting full nodes all the way up the tree

SPLIT-CHILD

- Inputs are a **nonfull** internal node **x** and an index **i** such that **y = x.c[i]** is a **full** child of **x**
 - Both nodes are assumed to be in main memory
- The procedure splits child **y** in two and adjusts **x** so that it has an additional child **z**
 - Node **y** is split about its median key, which is moved up into **y**'s parent node **x**
 - Keys in **y** that are greater than the median key are moved into a new node **z**, which becomes a new child of **x**


SPLIT-CHILD(x, i)

```
z := ALLOCATE-NODE()
y := x.c[i]
z.leaf := y.leaf
z.n := t - 1
for j := 0 to t - 2
    z.key[j] := y.key[j + t]
if not y.leaf
    for j := 0 to t - 1
        z.c[j] := y.c[j + t]
y.n := t - 1
for j := x.n downto i + 1
    x.c[j + 1] := x.c[j]
x.c[i + 1] := z
for j := x.n - 1 downto i
    x.key[j + 1] := x.key[j]
x.key[i] := y.key[t]
x.n := x.n + 1
DISK-WRITE(y)
DISK-WRITE(z)
DISK-WRITE(x)
```

Example

- **Splitting a node with $t = 2$**

- Each node can store at most $2t - 1 = 3$ keys
- Other children are omitted in the representation



SPLIT-CHILD(x, i)

```
z := ALLOCATE-NODE()
y := x.c[i]
z.leaf := y.leaf
z.n := t - 1
for j := 0 to t - 2
    z.key[j] := y.key[j + t]
if not y.leaf
    for j := 0 to t - 1
        z.c[j] := y.c[j + t]
y.n := t - 1
for j := x.n downto i + 1
    x.c[j + 1] := x.c[j]
x.c[i + 1] := z
for j := x.n - 1 downto i
    x.key[j + 1] := x.key[j]
x.key[i] := y.key[t]
x.n := x.n + 1
DISK-WRITE(y)
DISK-WRITE(z)
DISK-WRITE(x)
```

Analysis

- Total CPU time is $O(t)$
- $O(1)$ disk operations performed
- After **SPLIT-CHILD** is executed, the tree grows in height by one
 - Splitting is the only means by which the tree grows

```
SPLIT-CHILD(x,i)
  z := ALLOCATE-NODE()
  y := x.c[i]
  z.leaf := y.leaf
  z.n := t - 1
  for j := 0 to t - 2
    z.key[j] := y.key[j + t]
  if not y.leaf
    for j := 0 to t - 1
      z.c[j] := y.c[j + t]
  y.n := t - 1
  for j := x.n downto i + 1
    x.c[j + 1] := x.c[j]
  x.c[i + 1] := z
  for j := x.n - 1 downto i
    x.key[j + 1] := x.key[j]
  x.key[i] := y.key[t]
  x.n := x.n + 1
  DISK-WRITE(y)
  DISK-WRITE(z)
  DISK-WRITE(x)
```

INSERT-NONFULL

- **Procedure to insert key k into node x**
 - x is assumed to be nonfull when the procedure is called
- **There are two cases**
 1. If x is a leaf, insert key k in the appropriate position
 2. If x is not a leaf node, then insert k into the appropriate leaf node in the subtree rooted at node x

```
INSERT-NONFULL( $x, k$ )
 $i := x.n - 1$ 
if  $x.leaf$ 
    while  $i \geq 0$  and  $k < x.key[i]$ 
         $x.key[i + 1] := x.key[i]$ 
         $i := i - 1$ 
     $x.key[i + 1] := k$ 
     $x.n := x.n + 1$ 
    DISK-WRITE( $x$ )
else while  $i \geq 0$  and  $k < x.key[i]$ 
     $i := i - 1$ 
     $i := i + 1$ 
    DISK-READ( $x.c[i]$ )
    if  $x.c[i].n = 2t - 1$ 
        SPLIT-CHILD( $x, i$ )
        if  $k > x.key[i]$ 
             $i := i + 1$ 
    INSERT-NONFULL( $x.c[i], k$ )
```

INSERT-NONFULL (case 2)

- The while loop determines the child of **x** to which the recursion descends
- If recursion descends to a full child
 - Call **SPLIT-CHILD**
 - Then determine which of the two children is the correct one to descend to
- **One-pass strategy**
 - Every time we encounter a full node, we split it

```
INSERT-NONFULL(x,k)
  i := x.n - 1
  if x.leaf
    while i ≥ 0 and k < x.key[i]
      x.key[i + 1] := x.key[i]
      i := i - 1
    x.key[i + 1] := k
    x.n := x.n + 1
    DISK-WRITE(x)
  else while i ≥ 0 and k < x.key[i]
    i := i - 1
  i := i + 1
  DISK-READ(x.c[i])
  if x.c[i].n = 2t - 1
    SPLIT-CHILD(x,i)
    if k > x.key[i]
      i := i + 1
  INSERT-NONFULL(x.c[i],k)
```


INSERT

- Procedure to insert key **k** into a B-tree **T**
- If the root is full, call **SPLIT-CHILD** before calling **INSERT-NONFULL**
- The total CPU time used is **$O(t \log_t n)$**
- **$O(\log_t n)$** disk accesses

```
INSERT (T,k)
  r := T.root
  if r.n = 2t - 1
    s := ALLOCATE-NODE()
    T.root := s
    s.leaf := FALSE
    s.n := 0
    s.c[0] := r
    SPLIT-CHILD(s,0)
    INSERT-NONFULL(s,k)
  else
    INSERT-NONFULL(r,k)
```

Example

- **Add the sequence of keys given below to an empty B-tree with $t = 2$**
 - 5,7,3,12,24,6,2,1,9,11,8,4,10,23,21,15,16,17
 - Each node can store at most $2t - 1 = 3$ keys

Example

- **Add the sequence of keys given below to an empty B-tree with $t = 2$**
 - 5,7,3,12,24,6,2,1,9,11,8,4,10,23,21,15,16,17
 - Each node can store at most $2t - 1 = 3$ keys

5


Example

- **Add the sequence of keys given below to an empty B-tree with $t = 2$**
 - **5, 7, 3, 12, 24, 6, 2, 1, 9, 11, 8, 4, 10, 23, 21, 15, 16, 17**
 - Each node can store at most $2t - 1 = 3$ keys



Example


- **Add the sequence of keys given below to an empty B-tree with $t = 2$**
 - **5, 7, 3**, 12, 24, 6, 2, 1, 9, 11, 8, 4, 10, 23, 21, 15, 16, 17
 - Each node can store at most $2t - 1 = 3$ keys



full node

Example


- Add the sequence of keys given below to an empty B-tree with $t = 2$
 - 5, 7, 3, 12, 24, 6, 2, 1, 9, 11, 8, 4, 10, 23, 21, 15, 16, 17
 - Each node can store at most $2t - 1 = 3$ keys



SPLIT


Example

- Add the sequence of keys given below to an empty B-tree with $t = 2$
 - 5, 7, 3, 12, 24, 6, 2, 1, 9, 11, 8, 4, 10, 23, 21, 15, 16, 17
 - Each node can store at most $2t - 1 = 3$ keys



Example


- Add the sequence of keys given below to an empty B-tree with $t = 2$
 - 5, 7, 3, 12, 24, 6, 2, 1, 9, 11, 8, 4, 10, 23, 21, 15, 16, 17
 - Each node can store at most $2t - 1 = 3$ keys



SPLIT


Example

- Add the sequence of keys given below to an empty B-tree with $t = 2$
 - 5, 7, 3, 12, 24, 6, 2, 1, 9, 11, 8, 4, 10, 23, 21, 15, 16, 17
 - Each node can store at most $2t - 1 = 3$ keys




Example

- **Add the sequence of keys given below to an empty B-tree with $t = 2$**
 - 5, 7, 3, 12, 24, 6, 2, 1, 9, 11, 8, 4, 10, 23, 21, 15, 16, 17
 - Each node can store at most $2t - 1 = 3$ keys




Example

- **Add the sequence of keys given below to an empty B-tree with $t = 2$**
 - 5, 7, 3, 12, 24, 6, 2, 1, 9, 11, 8, 4, 10, 23, 21, 15, 16, 17
 - Each node can store at most $2t - 1 = 3$ keys




Example

- Add the sequence of keys given below to an empty B-tree with $t = 2$
 - 5, 7, 3, 12, 24, 6, 2, 1, 9, 11, 8, 4, 10, 23, 21, 15, 16, 17
 - Each node can store at most $2t - 1 = 3$ keys




Example

- Add the sequence of keys given below to an empty B-tree with $t = 2$
 - 5, 7, 3, 12, 24, 6, 2, 1, 9, 11, 8, 4, 10, 23, 21, 15, 16, 17
 - Each node can store at most $2t - 1 = 3$ keys




Example

- Add the sequence of keys given below to an empty B-tree with $t = 2$
 - 5, 7, 3, 12, 24, 6, 2, 1, 9, 11, 8, 4, 10, 23, 21, 15, 16, 17
 - Each node can store at most $2t - 1 = 3$ keys



Example


- Add the sequence of keys given below to an empty B-tree with $t = 2$
 - 5, 7, 3, 12, 24, 6, 2, 1, 9, 11, 8, 4, 10, 23, 21, 15, 16, 17
 - Each node can store at most $2t - 1 = 3$ keys



SPLIT


Example

- Add the sequence of keys given below to an empty B-tree with $t = 2$
 - 5, 7, 3, 12, 24, 6, 2, 1, 9, 11, 8, 4, 10, 23, 21, 15, 16, 17
 - Each node can store at most $2t - 1 = 3$ keys




Example

- Add the sequence of keys given below to an empty B-tree with $t = 2$
 - 5, 7, 3, 12, 24, 6, 2, 1, 9, 11, 8, 4, 10, 23, 21, 15, 16, 17
 - Each node can store at most $2t - 1 = 3$ keys




Example

- Add the sequence of keys given below to an empty B-tree with $t = 2$
 - 5, 7, 3, 12, 24, 6, 2, 1, 9, 11, 8, 4, 10, 23, 21, 15, 16, 17
 - Each node can store at most $2t - 1 = 3$ keys




Example

- Add the sequence of keys given below to an empty B-tree with $t = 2$
 - 5, 7, 3, 12, 24, 6, 2, 1, 9, 11, 8, 4, 10, 23, 21, 15, 16, 17
 - Each node can store at most $2t - 1 = 3$ keys



Example

- Add the sequence of keys given below to an empty B-tree with $t = 2$
 - 5, 7, 3, 12, 24, 6, 2, 1, 9, 11, 8, 4, 10, 23, 21, 15, 16, 17
 - Each node can store at most $2t - 1 = 3$ keys



Variants

- **2-3 trees** are B-trees in which every internal node has either two or three children
 - 2-3-4 trees can have up to 4 children (see slide 6)
- **B+trees** store all the satellite information in the leaves and stores only keys and child pointers in the internal nodes
 - To maximise their branching factor
- **B* trees** balance more neighbouring internal nodes to keep them more densely packed
 - This variant ensures non-root nodes are at least $\frac{2}{3}$ full instead of $\frac{1}{2}$

Applications

- **Example databases using B-trees**
 - MySQL uses both B-Tree and B+Tree indices
 - Oracle Database uses a B-tree index
 - Microsoft's SQL Server uses a B+Tree index
- **Example filesystems based on B-trees**
 - Apple's HFS+
 - Microsoft's NTFS
 - Linux Ext4
 - Reiser4 uses B*-trees
 - DragonFly BSD's HAMMER uses B+trees

Summary

- **Definition**
- **Motivation**
- **Properties**
- **Operations**
 - Search
 - Insertions
- **Variants**
- **Applications**